

Assignment 6

Link to questions - [here](#)

Student Details:

- Name : Anjishnu Mukherjee
- Registration Number : B05-511017020
- Class Roll Number : CS Gy-70
- Exam Roll Number : 510517086
- Email : 511017020.anjishnu@students.iiests.ac.in

Project Setup

Mount Google Drive

Mounted at /content/gdrive

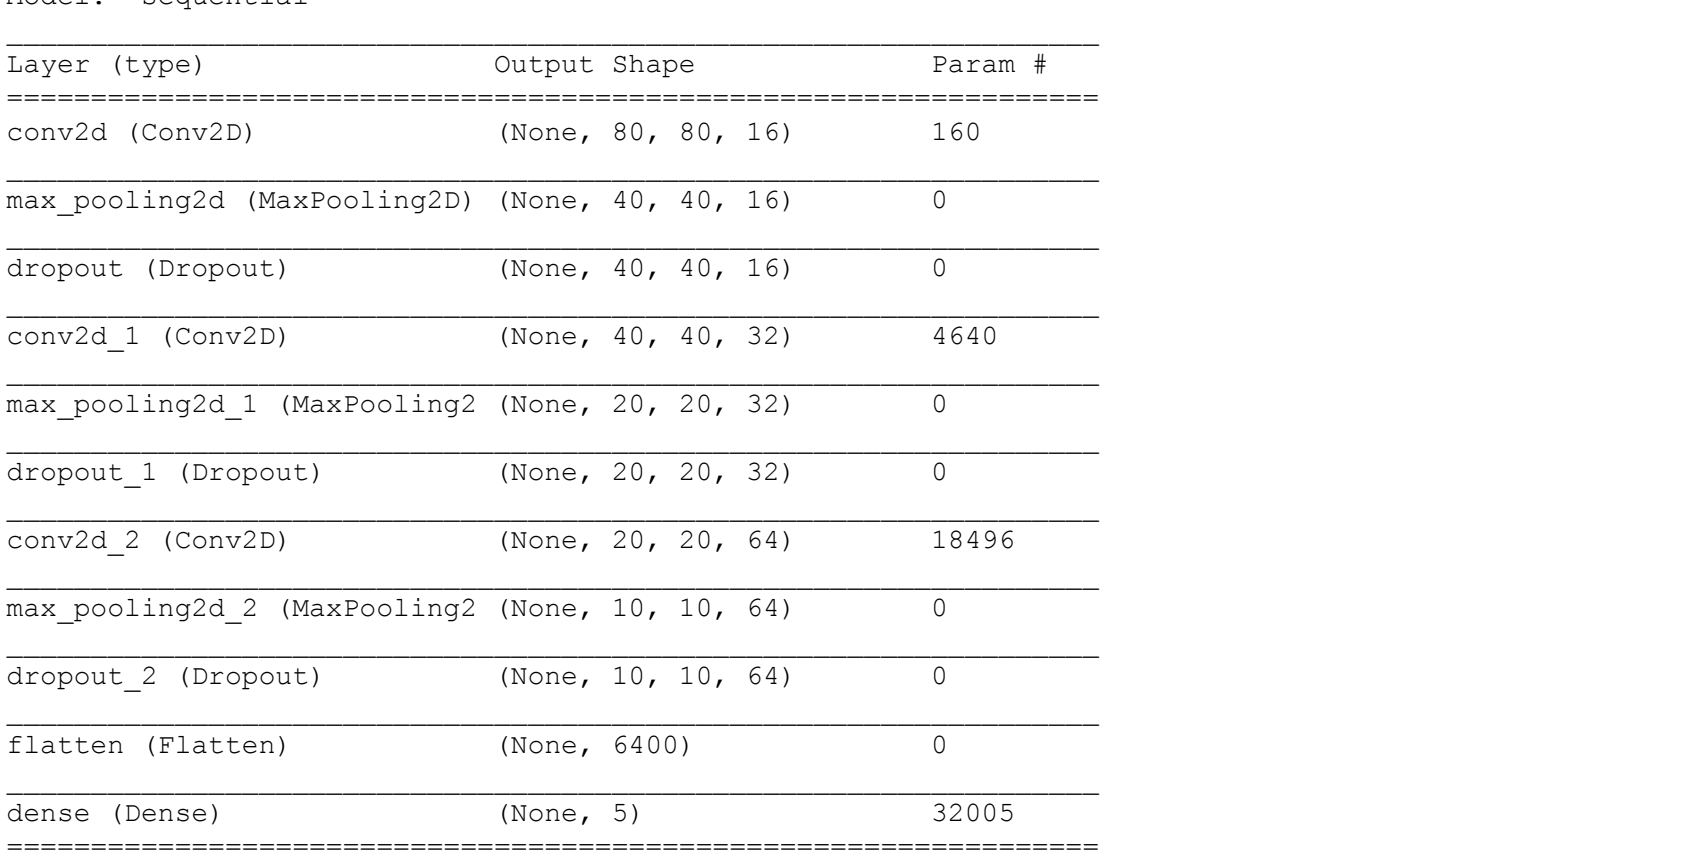
Importing all necessary libraries

Data-Loading

Create training data

Define helper functions for plotting accuracy, loss, confusion matrix for all models

Visualizing some of the images we loaded



Pre-processing

One-Hot Encoding

Train : Test split

Trying out the initial 4 models

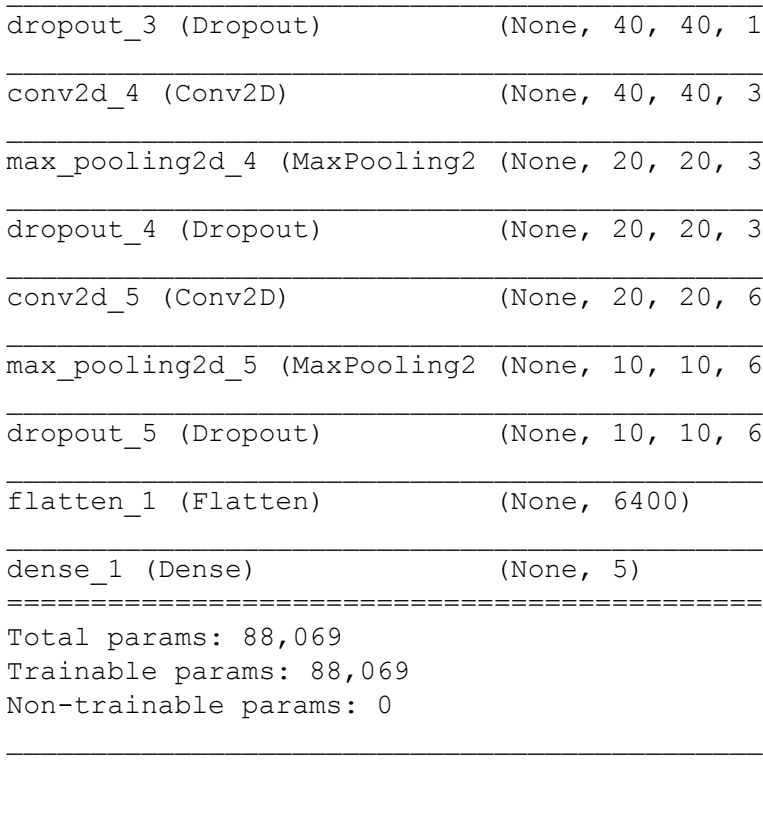
Model 1

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_1 (MaxPooling2D)	(None, 40, 40, 16)	0
dropout_1 (Dropout)	(None, 40, 40, 16)	0
conv2d_2_1 (Conv2D)	(None, 40, 40, 32)	4640
max_pooling2d_2_1 (MaxPooling2D)	(None, 20, 20, 32)	0
dropout_2_1 (Dropout)	(None, 20, 20, 32)	0
conv2d_2_2 (Conv2D)	(None, 20, 20, 64)	18496
max_pooling2d_2_2 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_2_2 (Dropout)	(None, 10, 10, 64)	0
flatten_1 (Flatten)	(None, 6400)	0
dense_1 (Dense)	(None, 5)	32005

=====
Total params: 55,301
Trainable params: 55,301
Non-trainable params: 0

Maximum training accuracy: 0.8999999761581421
Maximum validation accuracy: 0.5628019571304321



Confusion Matrix (Test samples)

	0	1	2	3	4
0	36,000	21,000	6,000	13,000	6,000
1	11,000	58,000	7,000	11,000	5,000
2	6,000	17,000	35,000	35,000	37,000
3	8,000	11,000	6,000	41,000	8,000
4	5,000	9,000	15,000	11,000	60,000

	precision	recall	f1-score	support
0	0.53	0.44	0.48	82
1	0.40	0.51	0.45	74
2	0.45	0.33	0.38	84
3	0.45	0.55	0.49	74
4	0.62	0.60	0.61	100

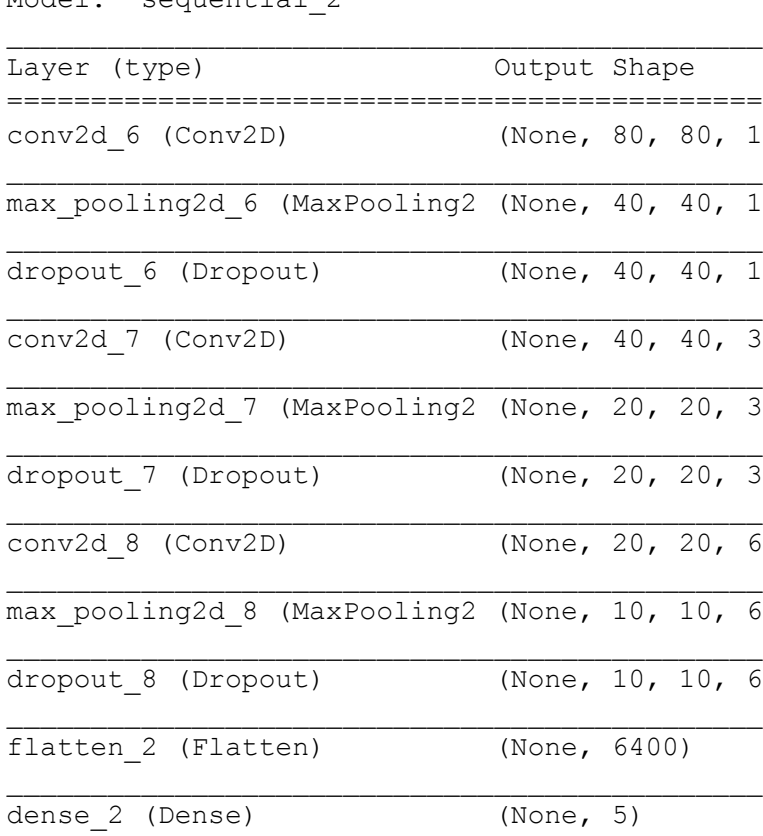
accuracy 0.49 0.49 0.48 0.49 414
macro avg 0.49 0.49 0.48 0.49 414
weighted avg 0.50 0.49 0.49 0.49 414

Model 2

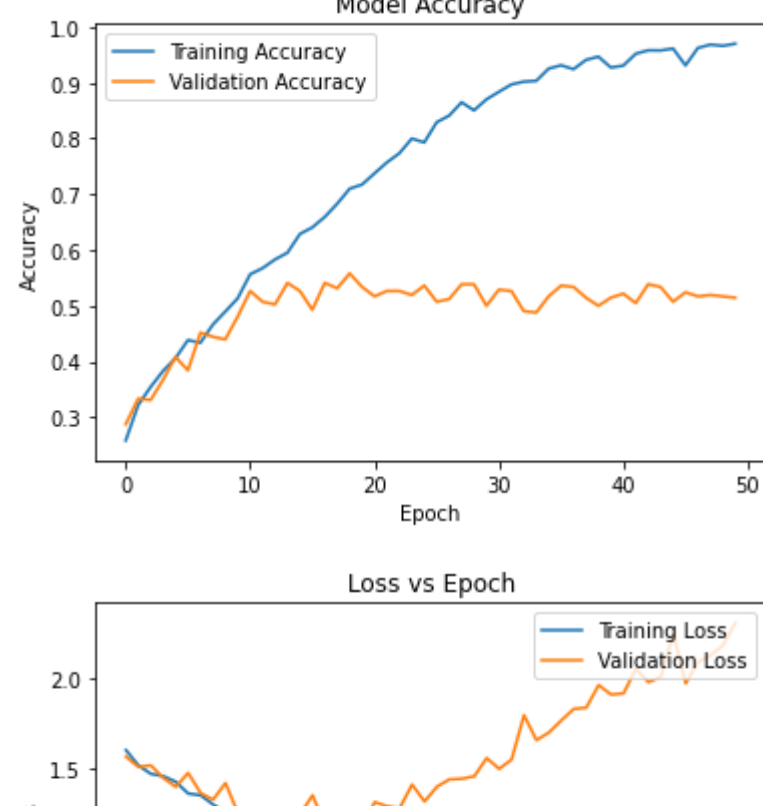
Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_3 (MaxPooling2D)	(None, 40, 40, 16)	0
dropout_3 (Dropout)	(None, 40, 40, 16)	0
conv2d_4 (Conv2D)	(None, 40, 40, 32)	4640
max_pooling2d_4 (MaxPooling2D)	(None, 20, 20, 32)	0
dropout_4 (Dropout)	(None, 20, 20, 32)	0
conv2d_5 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_5 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_5 (Dropout)	(None, 10, 10, 64)	0
flatten_1 (Flatten)	(None, 6400)	0
dense_1 (Dense)	(None, 5)	32005

=====
Total params: 89,069
Trainable params: 89,069
Non-trainable params: 0



Maximum training accuracy: 0.9620967507362366
Maximum validation accuracy: 0.555555820465088



	precision	recall	f1-score	support
0	0.58	0.50	0.54	82
1	0.46	0.50	0.48	74
2	0.39	0.35	0.37	84
3	0.42	0.47	0.45	74
4	0.54	0.57	0.55	100

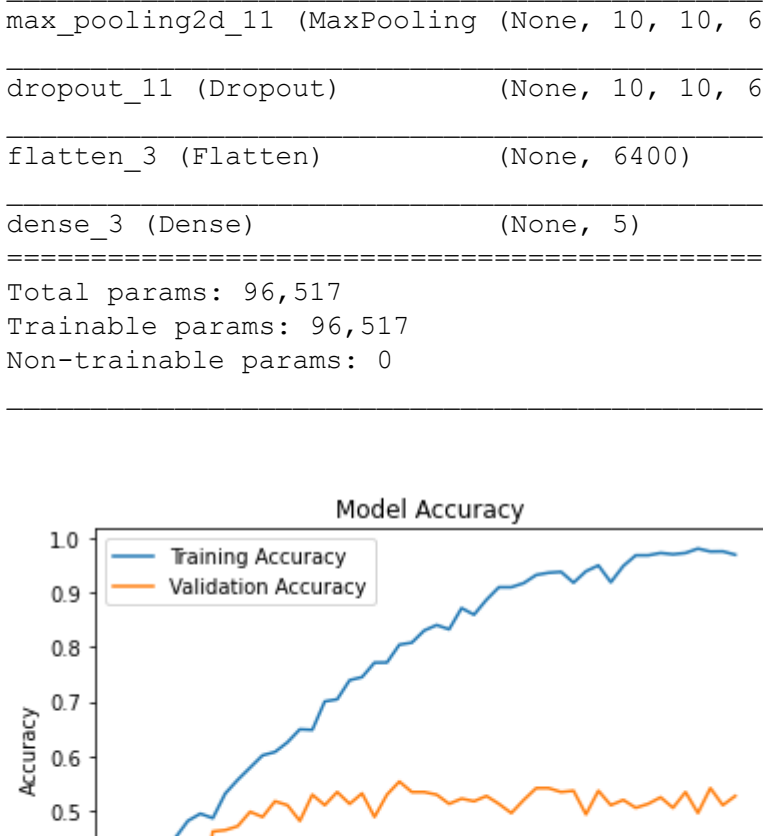
accuracy 0.48 0.48 0.48 0.48 414
macro avg 0.48 0.48 0.48 0.48 414
weighted avg 0.48 0.48 0.48 0.48 414

Model 3


Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_6 (MaxPooling2D)	(None, 40, 40, 16)	0
dropout_6 (Dropout)	(None, 40, 40, 16)	0
conv2d_7 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_7 (MaxPooling2D)	(None, 20, 20, 32)	0
dropout_7 (Dropout)	(None, 20, 20, 32)	0
conv2d_8 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_8 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_8 (Dropout)	(None, 10, 10, 64)	0
flatten_2 (Flatten)	(None, 6400)	0
dense_2 (Dense)	(None, 5)	32005

=====
Total params: 96,261
Trainable params: 96,261
Non-trainable params: 0



Maximum training accuracy: 0.9701613187789917
Maximum validation accuracy: 0.557971006713867



	precision	recall	f1-score	support
0	0.54	0.49	0.51	82
1	0.44	0.64	0.52	74
2	0.50	0.26	0.34	84
3	0.49	0.64	0.55	74
4	0.61	0.57	0.59	100

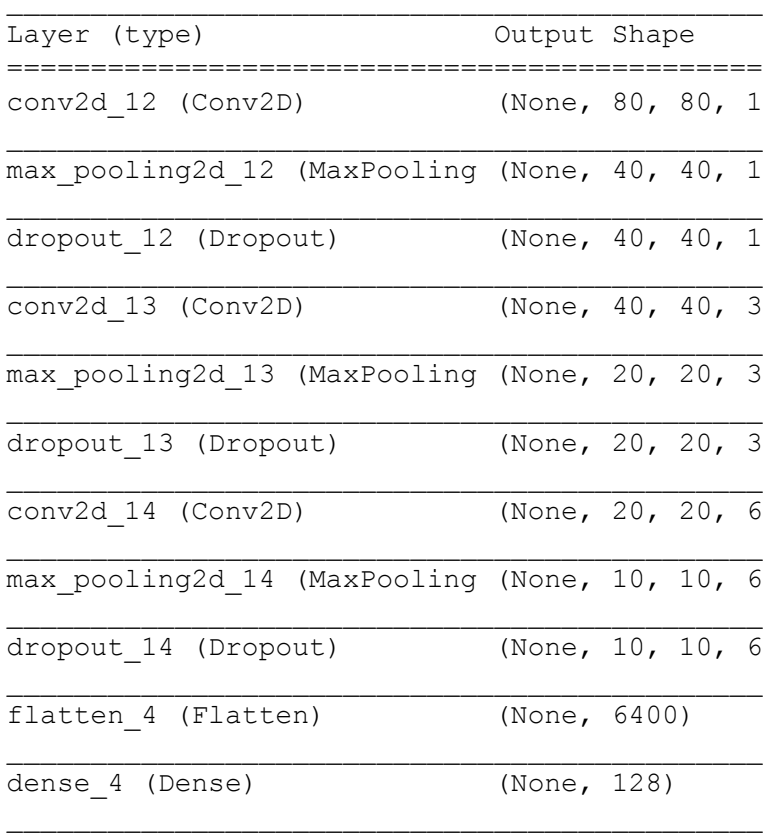
accuracy 0.52 0.52 0.51 0.51 414
macro avg 0.52 0.52 0.50 0.50 414
weighted avg 0.52 0.51 0.51 0.51 414

Model 4

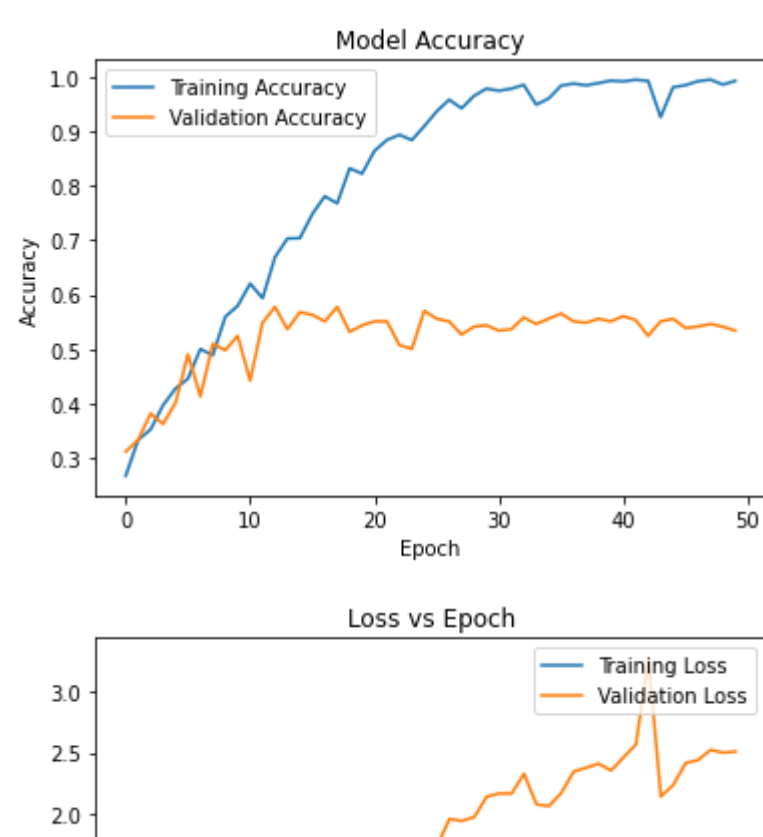
Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 80, 80, 16)	416
max_pooling2d_9 (MaxPooling2D)	(None, 40, 40, 16)	0
dropout_9 (Dropout)	(None, 40, 40, 16)	0
conv2d_10 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_10 (MaxPooling2D)	(None, 20, 20, 32)	0
dropout_10 (Dropout)	(None, 20, 20, 32)	0
conv2d_11 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_11 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_11 (Dropout)	(None, 10, 10, 64)	0
flatten_3 (Flatten)	(None, 6400)	0
dense_3 (Dense)	(None, 5)	32005

=====
Total params: 96,517
Trainable params: 96,517
Non-trainable params: 0



Maximum training accuracy: 0.9814516305923462
Maximum validation accuracy: 0.5331401038169861



	precision	recall	f1-score	support
0	0.63	0.46	0.54	82
1	0.49	0.64	0.55	74
2	0.48	0.40	0.44	84
3	0.59	0.59	0.59	74
4	0.54	0.62	0.58	100

accuracy 0.53 0.53 0.52 0.52 414
macro avg 0.53 0.53 0.52 0.52 414
weighted avg 0.53 0.53 0.53 0.53 414

Model 3 was the best model with an accuracy of 57.84% on the validation dataset.

The model is as follows:

```
model3 = Sequential()
model3.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = 'relu', input_shape = (8, 80,1)))
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))
model3.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = 'relu'))
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))
model3.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = 'relu'))
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))
model3.add(Flatten())
model3.add(Dense(5, activation = 'softmax'))
model3.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model3.summary()
```

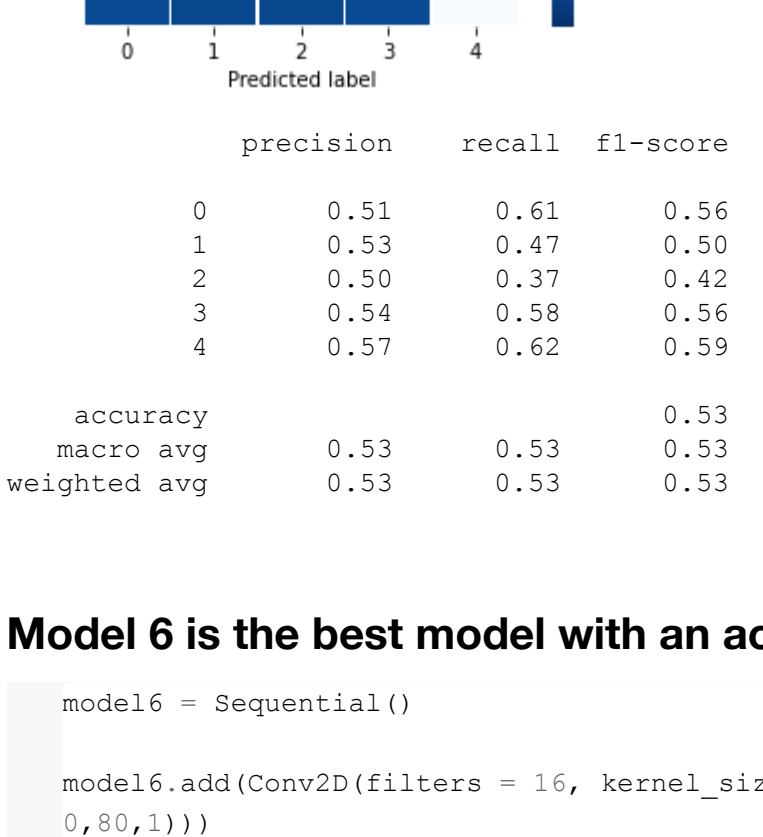
Now, we will be adding 2/3 fully-connected layers to this model.

Model 5

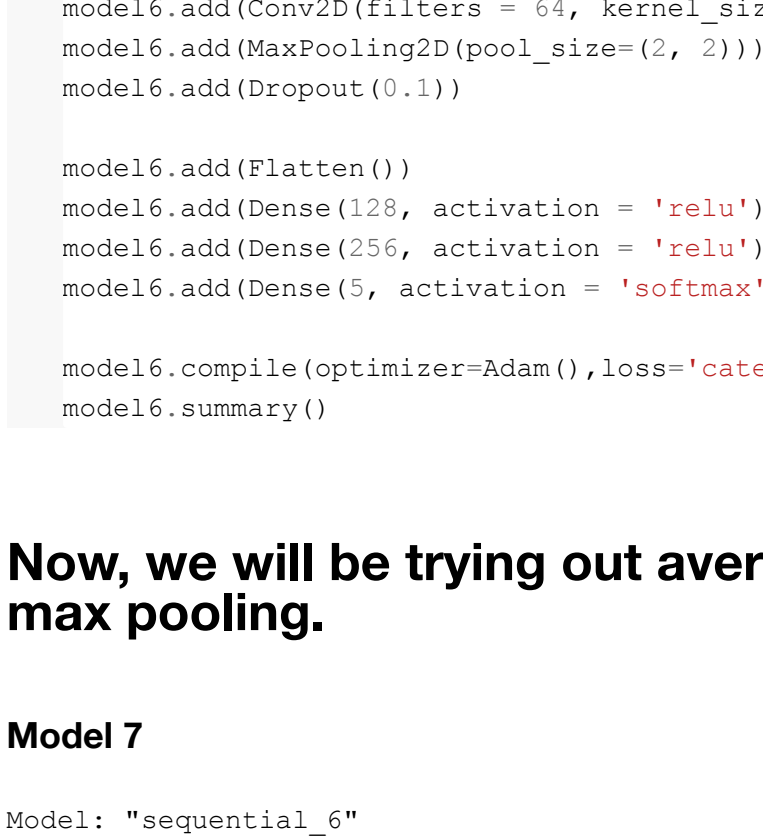
Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_12 (MaxPooling2D)	(None, 40, 40, 16)	0
dropout_12 (Dropout)	(None, 40, 40, 16)	0
conv2d_13 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_13 (MaxPooling2D)	(None, 20, 20, 32)	0
dropout_13 (Dropout)	(None, 20, 20, 32)	0
conv2d_14 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_14 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_14 (Dropout)	(None, 10, 10, 64)	0
flatten_4 (Flatten)	(None, 6400)	0
dense_4 (Dense)	(None, 128)	819328
dense_5 (Dense)	(None, 5)	645

=====
Total params: 884,229
Trainable params: 884,229
Non-trainable params: 0



Maximum training accuracy: 0.9946236610412598
Maximum validation accuracy: 0.5772947072982788



	precision	recall	f1-score	support
0	0.45	0.60	0.51	82
1	0.49	0.46	0.48	74
2	0.48	0.40	0.44	84
3	0.59	0.59	0.59	74
4	0.67	0.60	0.63	100

accuracy 0.54 0.53 0.53 0.53 414
macro avg 0.54 0.53 0.53 0.53 414
weighted avg 0.54 0.53 0.53 0.53 414

Model 6

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_15 (MaxPooling2D)	(None, 40, 40, 16)	0
dropout_15 (Dropout)	(None, 40, 40, 16)	0
conv2d_16 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_16 (MaxPooling2D)	(None, 20, 20, 32)	0
dropout_16 (Dropout)	(None, 20, 20, 32)	0
conv2d_17 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_17 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_17 (Dropout)	(None, 10, 10, 64)	0
flatten_5 (Flatten)	(None, 6400)	0
dense_6 (Dense)	(None, 128)	819328
dense_7 (Dense)	(None, 256)	33024
dense_8 (Dense)	(None, 5)	1285

=====
Total params: 917,893
Trainable params: 917,893
Non-trainable params: 0

Maximum training accuracy: 0.9975806474685669
Maximum validation accuracy: 0.5531401038169861

	precision	recall	f1-score	support
0	0.51	0.61	0.56	82
1	0.53	0.47	0.50	74
2	0.50	0.37	0.42	84
3	0.54	0.58	0.56	74
4	0.57	0.62	0.59	100

accuracy 0.53 0.53 0.53 0.53 414
macro avg 0.53 0.53 0.53 0.53 414
weighted avg 0.53 0.53 0.53 0.53 414

Model 6 is the best model with an accuracy of 58.82% on the validation dataset.

```
model6 = Sequential()
model6.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = 'relu', input_shape = (8, 80,1)))
model6.add(MaxPooling2D(pool_size=(2, 2)))
model6.add(Dropout(0.1))
model6.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = 'relu'))
model6.add(MaxPooling2D(pool_size=(2, 2)))
model6.add(Dropout(0.1))
model6.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = 'relu'))
model6.add(MaxPooling2D(pool_size=(2, 2)))
model6.add(Dropout(0.1))
model6.add(Flatten())
model6.add(Dense(128, activation = 'relu'))
model6.add(Dense(256, activation = 'relu'))
model6.add(Dense(5, activation = 'softmax'))
model6.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model6.summary()
```

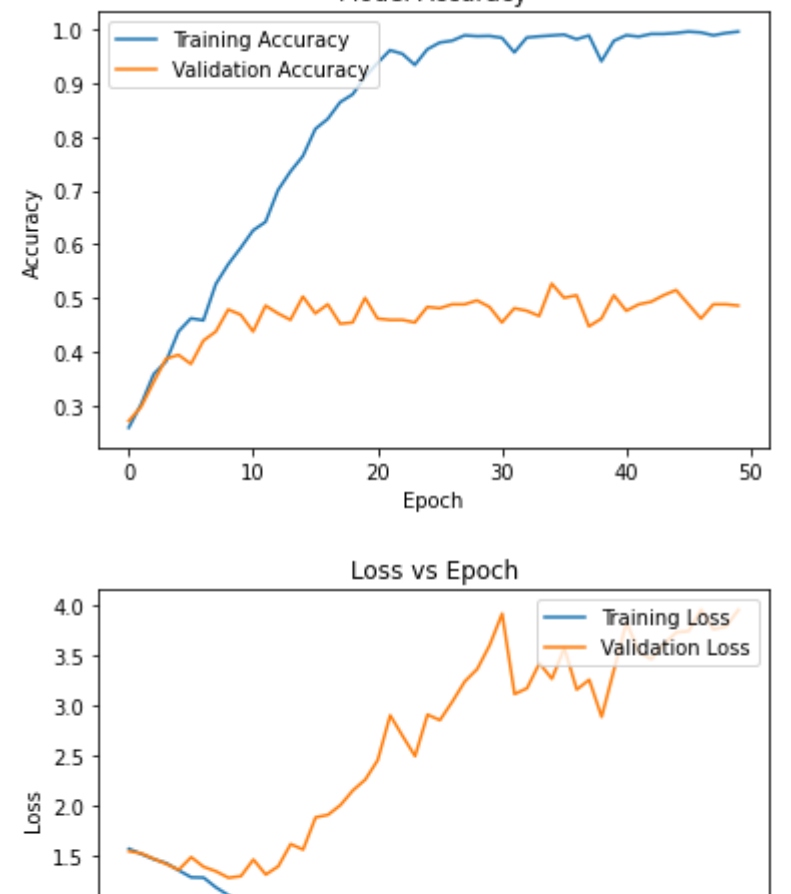
Now, we will be trying out average pooling on the same model instead of max pooling.

Model 7

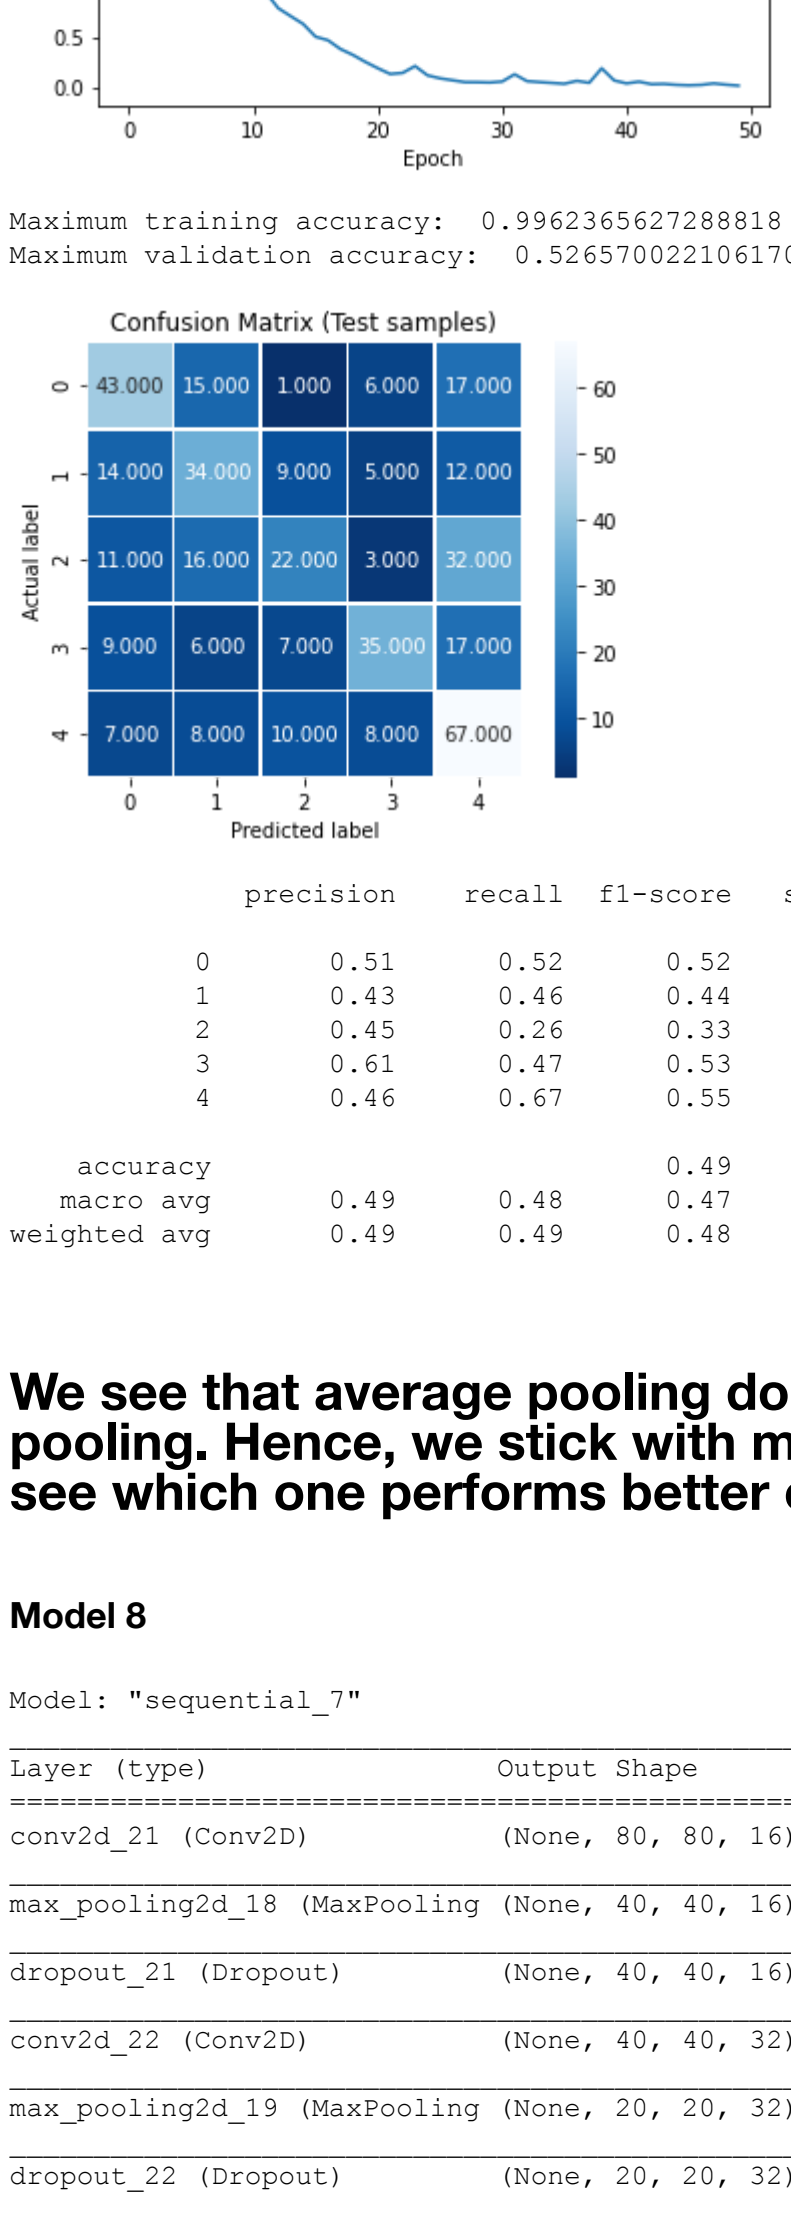
Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 80, 80, 16)	160
average_pooling2d_18 (AveragePooling2D)	(None, 40, 40, 16)	0
dropout_18 (Dropout)	(None, 40, 40, 16)	0
conv2d_19 (Conv2D)	(None, 40, 40, 32)	12832
average_pooling2d_19 (AveragePooling2D)	(None, 20, 20, 32)	0
dropout_19 (Dropout)	(None, 20, 20, 32)	0
conv2d_20 (Conv2D)	(None, 20, 20, 64)	51264
average_pooling2d_20 (AveragePooling2D)	(None, 10, 10, 64)	0
dropout_20 (Dropout)	(None, 10, 10, 64)	0
flatten_6 (Flatten)	(None, 6400)	0
dense_9 (Dense)	(None, 128)	819328
dense_10 (Dense)	(None, 256)	33024
dense_11 (Dense)	(None, 5)	1285

=====
Total params: 917,893
Trainable params: 917,893
Non-trainable params: 0



Maximum training accuracy: 0.9962365627288818
Maximum validation accuracy: 0.5265700221061707



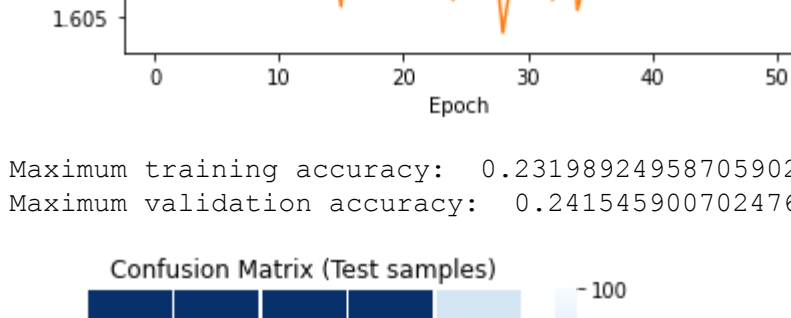
We see that average pooling does not give us better results than max pooling. Hence, we stick with model 6 and vary the activation functions to see which one performs better on model 6.

Model 8

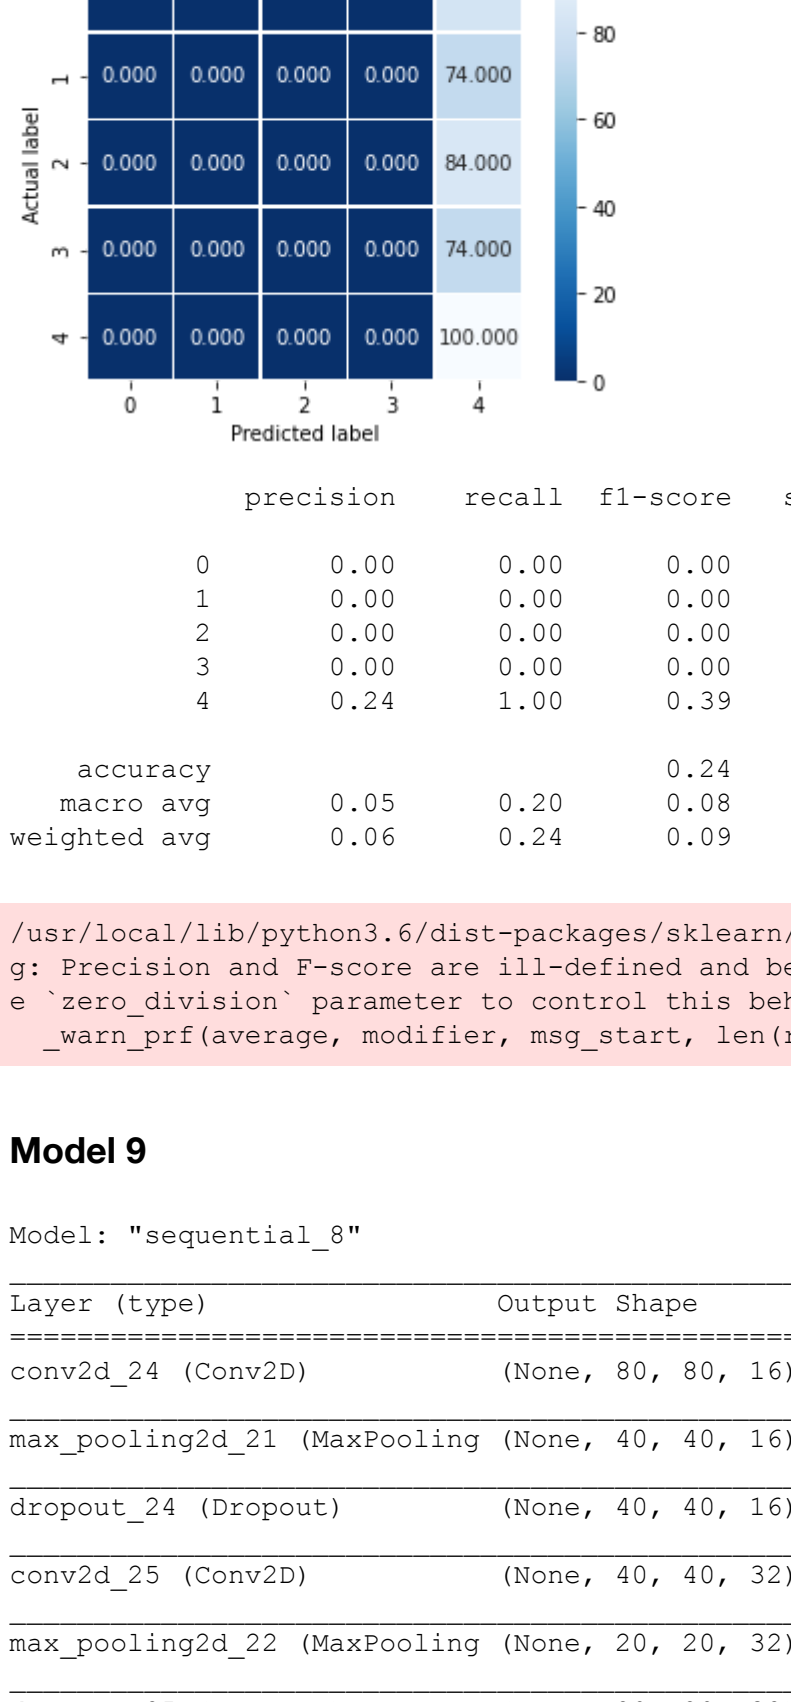
Model: "sequential_7"		
Layer (type)	Output Shape	Param #

conv2d_21 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_18 (MaxPooling)	(None, 40, 40, 16)	0
dropout_21 (Dropout)	(None, 40, 40, 16)	0
conv2d_22 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_19 (MaxPooling)	(None, 20, 20, 32)	0
dropout_22 (Dropout)	(None, 20, 20, 32)	0
conv2d_23 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_20 (MaxPooling)	(None, 10, 10, 64)	0
dropout_23 (Dropout)	(None, 10, 10, 64)	0
flatten_7 (Flatten)	(None, 6400)	0
dense_12 (Dense)	(None, 128)	819328
dense_13 (Dense)	(None, 256)	33024
dense_14 (Dense)	(None, 5)	1285

Total params: 917,893		
Trainable params: 917,893		
Non-trainable params: 0		



Maximum training accuracy: 0.23198924959705902
Maximum validation accuracy: 0.2415459007024765




/usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
warn_prf(average, modifier, msg_start, len(result))

Model 9

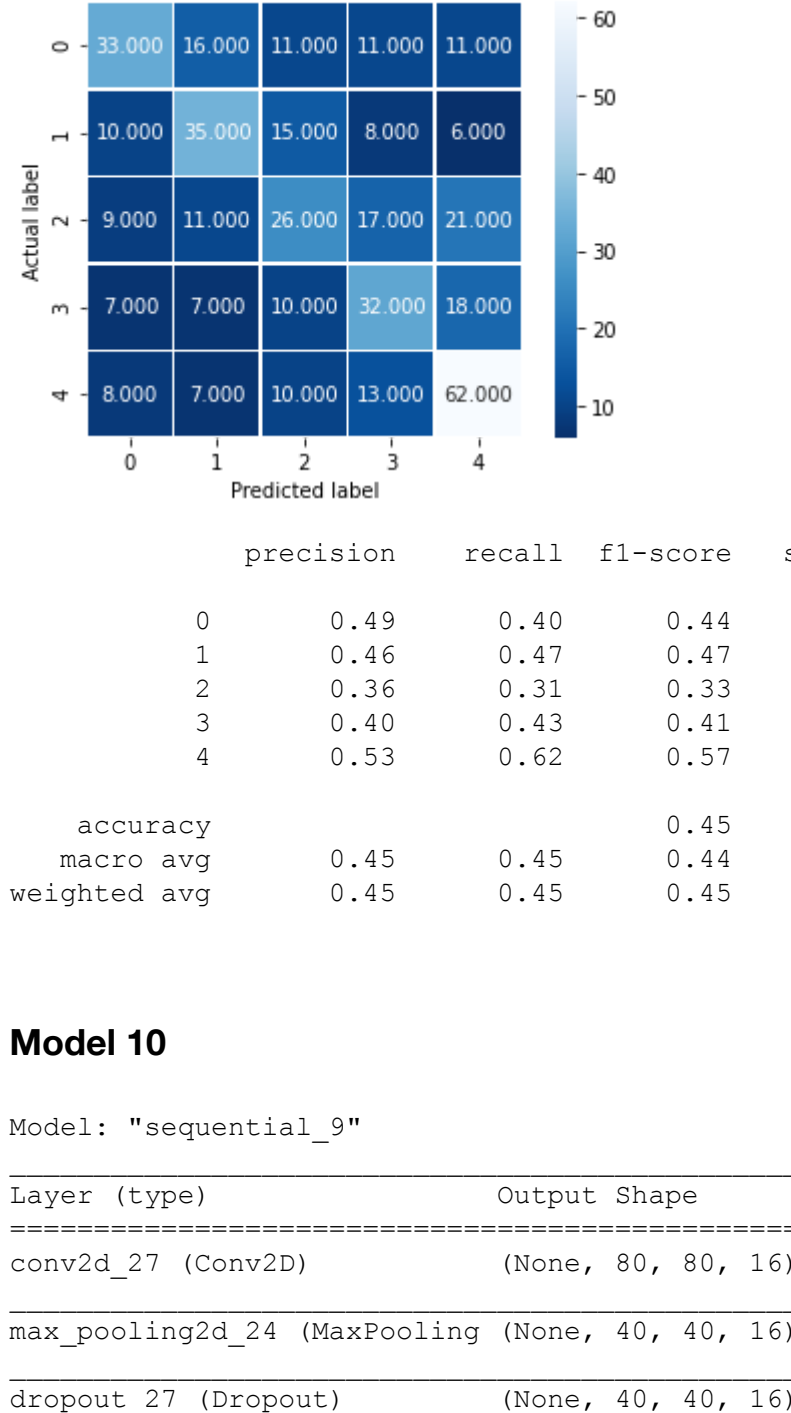
Model: "sequential_8"		
Layer (type)	Output Shape	Param #

conv2d_24 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_21 (MaxPooling)	(None, 40, 40, 16)	0
dropout_24 (Dropout)	(None, 40, 40, 16)	0
conv2d_25 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_22 (MaxPooling)	(None, 20, 20, 32)	0
dropout_25 (Dropout)	(None, 20, 20, 32)	0
conv2d_26 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_23 (MaxPooling)	(None, 10, 10, 64)	0
dropout_26 (Dropout)	(None, 10, 10, 64)	0
flatten_8 (Flatten)	(None, 6400)	0
dense_15 (Dense)	(None, 128)	819328
dense_16 (Dense)	(None, 256)	33024
dense_17 (Dense)	(None, 5)	1285

Total params: 917,893		
Trainable params: 917,893		
Non-trainable params: 0		



Maximum training accuracy: 0.9962365627288818
Maximum validation accuracy: 0.4830938325500485




Model 10


Model: "sequential_9"		
Layer (type)	Output Shape	Param #

conv2d_27 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_24 (MaxPooling)	(None, 40, 40, 16)	0
dropout_27 (Dropout)	(None, 40, 40, 16)	0
conv2d_28 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_25 (MaxPooling)	(None, 20, 20, 32)	0
dropout_28 (Dropout)	(None, 20, 20, 32)	0
conv2d_29 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_26 (MaxPooling)	(None, 10, 10, 64)	0
dropout_29 (Dropout)	(None, 10, 10, 64)	0
flatten_9 (Flatten)	(None, 6400)	0
dense_18 (Dense)	(None, 128)	819328
dense_19 (Dense)	(None, 256)	33024
dense_20 (Dense)	(None, 5)	1285

Total params: 917,893		
Trainable params: 917,893		
Non-trainable params: 0		



Maximum training accuracy: 0.998182813644409
Maximum validation accuracy: 0.5990338325500488



After modifying the activation functions, we see that model 10's accuracy is the highest.

```
leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model10 = Sequential()

model10.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (50,50,1)))
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))

model10.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))

model10.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))

model10.add(Flatten())
model10.add(Dense(128, activation = leakyrelu))
model10.add(Dense(256, activation = leakyrelu))
model10.add(Dense(5, activation = 'softmax'))

model10.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model10.summary()
```


Thus, we will now vary the regularization parameters for model 10.

Model 11


Model: "sequential_10"		
Layer (type)	Output Shape	Param #

conv2d_30 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_27 (MaxPooling)	(None, 40, 40, 16)	0
dropout_30 (Dropout)	(None, 40, 40, 16)	0
conv2d_31 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_28 (MaxPooling)	(None, 20, 20, 32)	0
dropout_31 (Dropout)	(None, 20, 20, 32)	0
conv2d_32 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_29 (MaxPooling)	(None, 10, 10, 64)	0
dropout_32 (Dropout)	(None, 10, 10, 64)	0
flatten_10 (Flatten)	(None, 6400)	0
dense_21 (Dense)	(None, 128)	819328
dense_22 (Dense)	(None, 256)	33024
dense_23 (Dense)	(None, 5)	1285

Total params: 917,893		
Trainable params: 917,893		
Non-trainable params: 0		



Maximum training accuracy: 0.9927419624057007
Maximum validation accuracy: 0.5990338325500488




Model 12

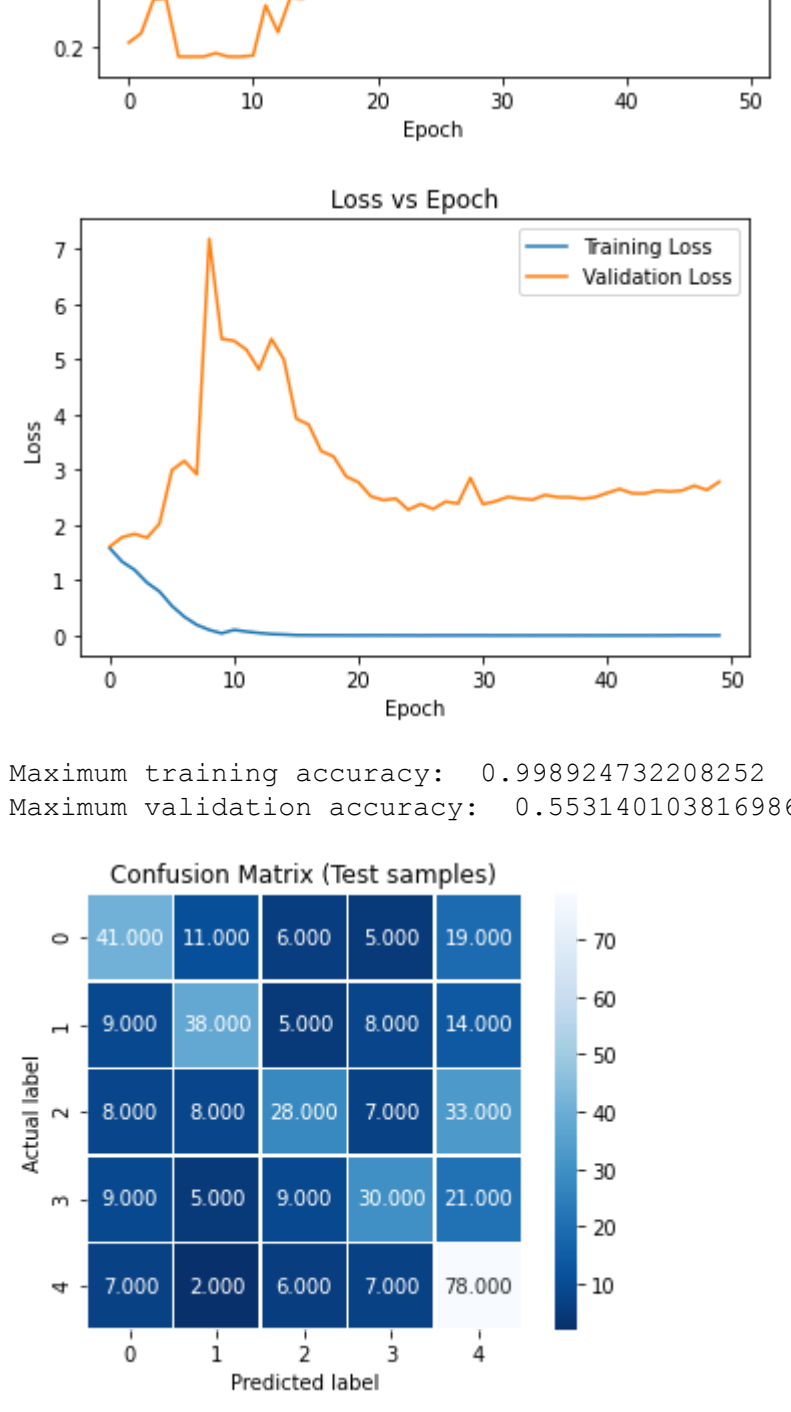
Model: "sequential_11"		
Layer (type)	Output Shape	Param #

conv2d_33 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_30 (MaxPooling)	(None, 40, 40, 16)	0
conv2d_34 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_31 (MaxPooling)	(None, 20, 20, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 20, 20, 32)	128
conv2d_35 (Conv2D)	(None, 20, 20, 64)	51264
batch_normalization_2 (Batch Normalization)	(None, 10, 10, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 10, 10, 64)	256
flatten_11 (Flatten)	(None, 6400)	0
dense_24 (Dense)	(None, 128)	819328
dense_25 (Dense)	(None, 256)	33024
dense_26 (Dense)	(None, 5)	1285

Total params: 918,277		
Trainable params: 918,085		
Non-trainable params: 192		



Maximum training accuracy: 0.998924732208202
Maximum validation accuracy: 0.5845410832500486

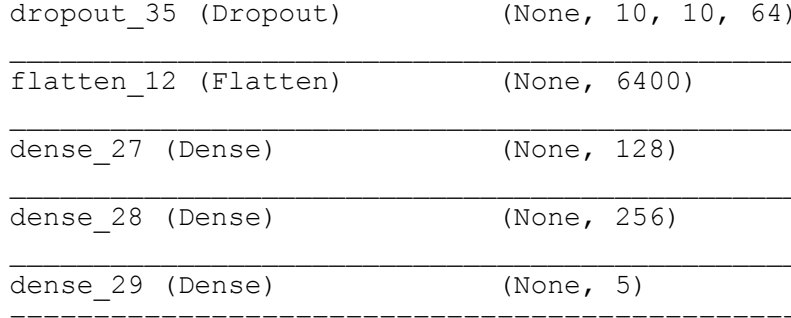


Model 13

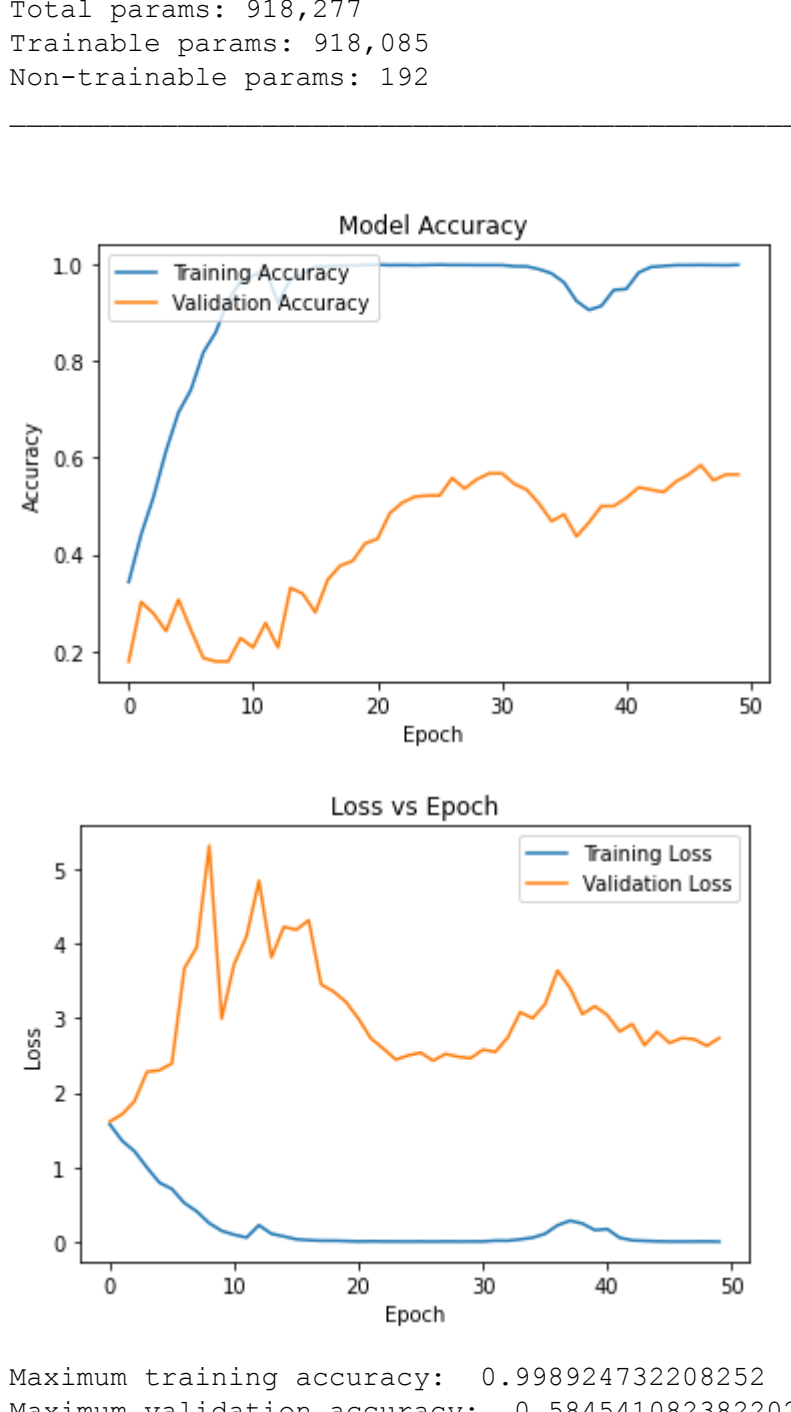
Model: "sequential_12"		
Layer (type)	Output Shape	Param #

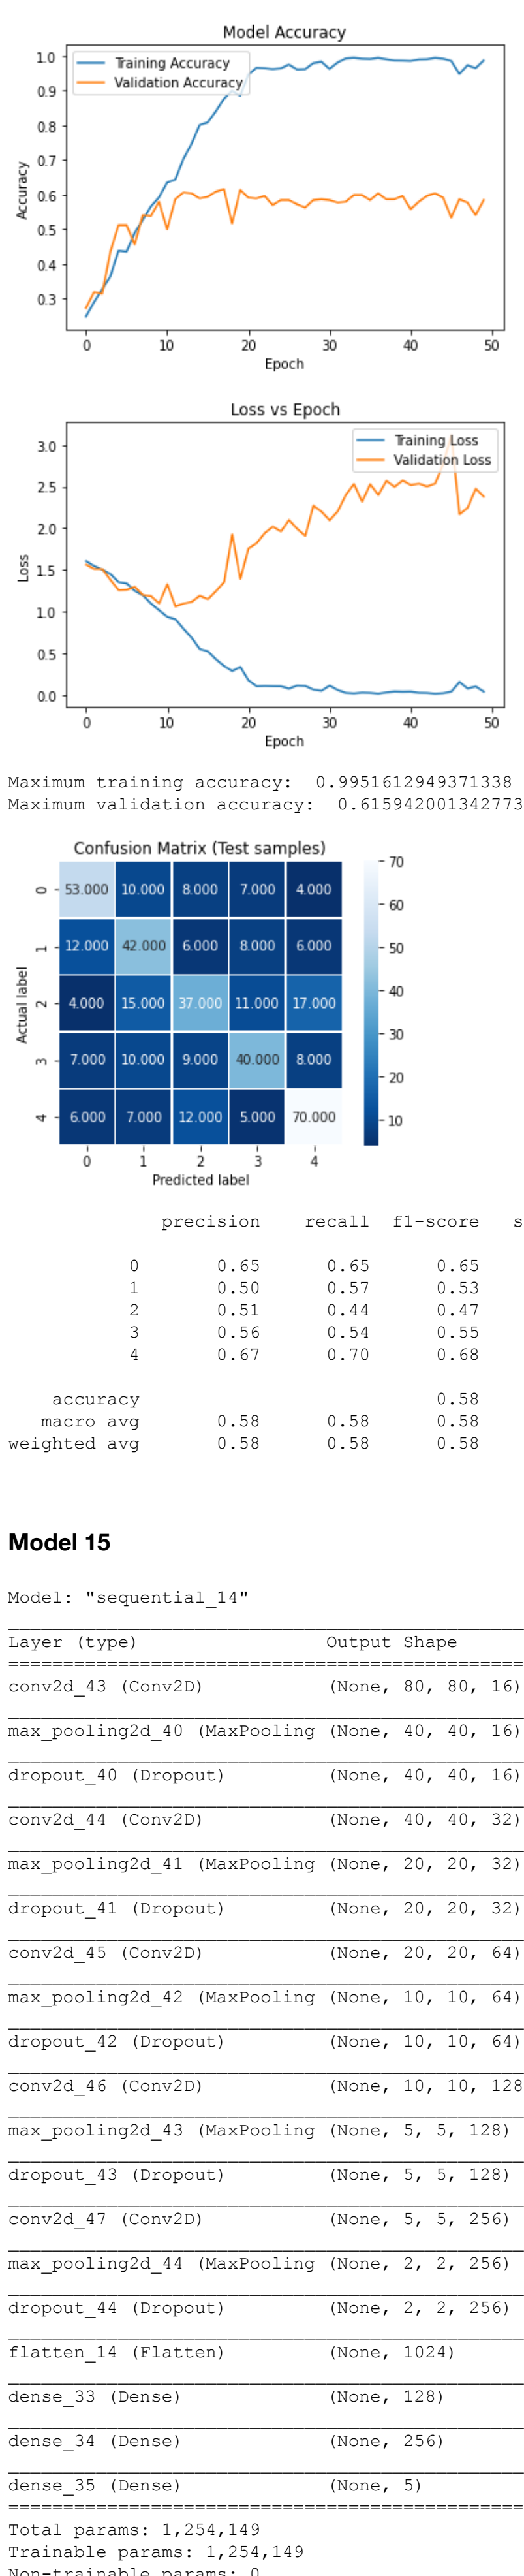
conv2d_36 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_33 (MaxPooling)	(None, 40, 40, 16)	0
dropout_33 (Dropout)	(None, 40, 40, 16)	0
conv2d_37 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_34 (MaxPooling)	(None, 20, 20, 32)	0
batch_normalization_2 (Batch Normalization)	(None, 20, 20, 32)	128
dropout_34 (Dropout)	(None, 20, 20, 32)	0
conv2d_38 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_35 (MaxPooling)	(None, 10, 10, 64)	0
batch_normalization_3 (Batch Normalization)	(None, 10, 10, 64)	256
dropout_35 (Dropout)	(None, 10, 10, 64)	0
flatten_12 (Flatten)	(None, 6400)	0
dense_27 (Dense)	(None, 128)	819328
dense_28 (Dense)	(None, 256)	33024
dense_29 (Dense)	(None, 5)	1285

Total params: 918,221		
Trainable params: 918,221		
Non-trainable params: 0		



Maximum training accuracy: 0.998924732208202
Maximum validation accuracy: 0.5845410832500486



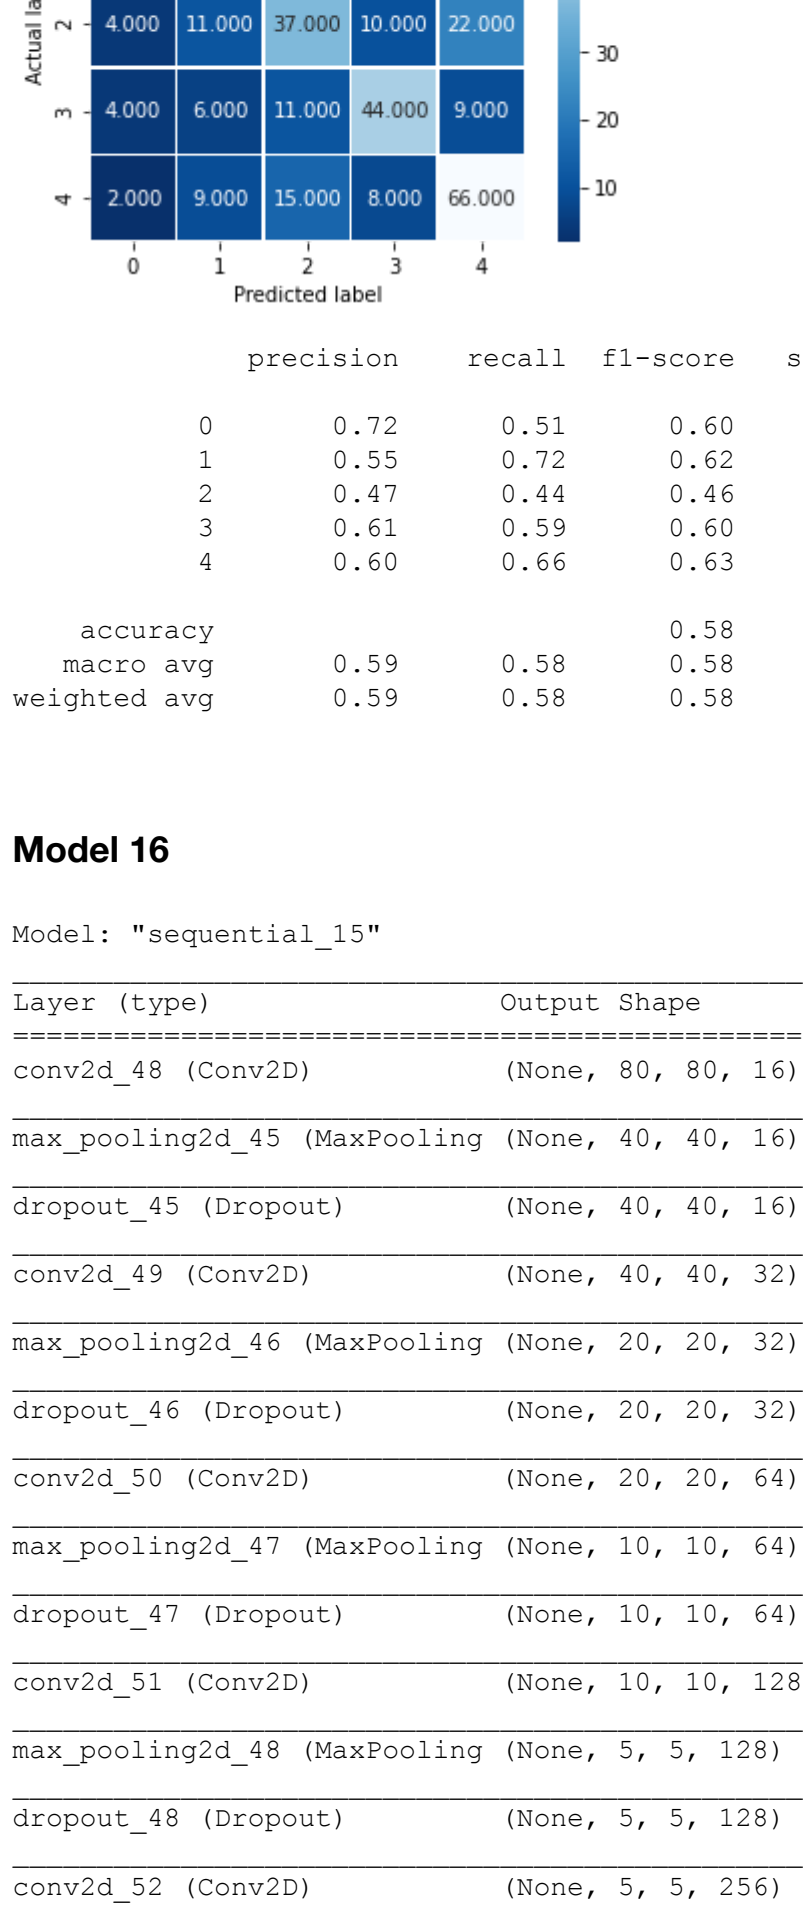


Model 15

Model: "sequential_14"

Layer (type)	Output Shape	Param #
conv2d_43 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_40 (MaxPooling)	(None, 40, 40, 16)	0
dropout_40 (Dropout)	(None, 40, 40, 16)	0
conv2d_44 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_41 (MaxPooling)	(None, 20, 20, 32)	0
dropout_41 (Dropout)	(None, 20, 20, 32)	0
conv2d_45 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_42 (MaxPooling)	(None, 10, 10, 64)	0
dropout_42 (Dropout)	(None, 10, 10, 64)	0
conv2d_46 (Conv2D)	(None, 10, 10, 128)	204928
max_pooling2d_43 (MaxPooling)	(None, 5, 5, 128)	0
dropout_43 (Dropout)	(None, 5, 5, 128)	0
conv2d_47 (Conv2D)	(None, 5, 5, 256)	819456
max_pooling2d_44 (MaxPooling)	(None, 2, 2, 256)	0
dropout_44 (Dropout)	(None, 2, 2, 256)	0
flatten_14 (Flatten)	(None, 1024)	0
dense_33 (Dense)	(None, 128)	131200
dense_34 (Dense)	(None, 256)	33024
dense_35 (Dense)	(None, 5)	1285

=====
Total params: 1,254,149
Trainable params: 1,254,149
Non-trainable params: 0

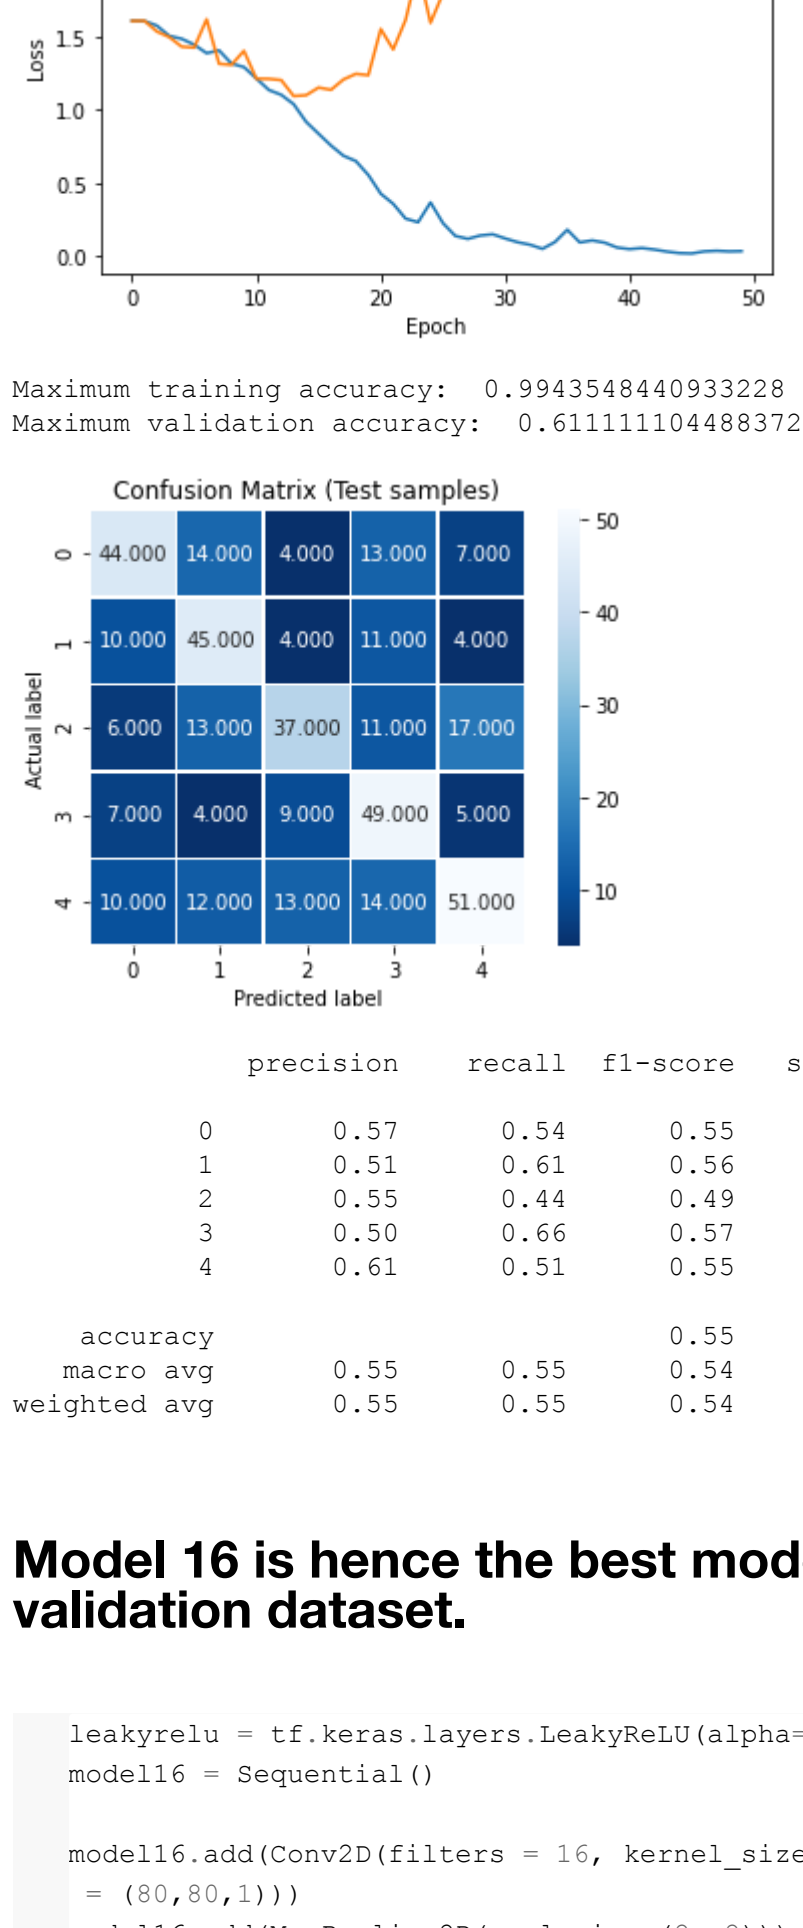


Model 16

Model: "sequential_15"

Layer (type)	Output Shape	Param #
conv2d_48 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_45 (MaxPooling)	(None, 40, 40, 16)	0
dropout_45 (Dropout)	(None, 40, 40, 16)	0
conv2d_49 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_46 (MaxPooling)	(None, 20, 20, 32)	0
dropout_46 (Dropout)	(None, 20, 20, 32)	0
conv2d_50 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_47 (MaxPooling)	(None, 10, 10, 64)	0
dropout_47 (Dropout)	(None, 10, 10, 64)	0
conv2d_51 (Conv2D)	(None, 10, 10, 128)	204928
max_pooling2d_48 (MaxPooling)	(None, 5, 5, 128)	0
dropout_48 (Dropout)	(None, 5, 5, 128)	0
conv2d_52 (Conv2D)	(None, 5, 5, 256)	819456
max_pooling2d_49 (MaxPooling)	(None, 2, 2, 256)	0
dropout_49 (Dropout)	(None, 2, 2, 256)	0
conv2d_53 (Conv2D)	(None, 2, 2, 512)	3277312
max_pooling2d_50 (MaxPooling)	(None, 1, 1, 512)	0
dropout_50 (Dropout)	(None, 1, 1, 512)	0
flatten_15 (Flatten)	(None, 512)	0
dense_36 (Dense)	(None, 128)	65664
dense_37 (Dense)	(None, 256)	33024
dense_38 (Dense)	(None, 5)	1285

=====
Total params: 4,465,825
Trainable params: 4,465,825
Non-trainable params: 0



Model 16 is hence the best model so far with an accuracy of 63.87% on the validation dataset.

```
leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model16 = Sequential()

model16.add(Conv2D(filters=16, kernel_size=(3,3),padding='Same',activation=leakyrelu,input_shape=(50,50,1)))
model16.add(MaxPooling2D(pool_size=(2,2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters=32, kernel_size=(5,5),padding='Same',activation=leakyrelu))
model16.add(MaxPooling2D(pool_size=(2,2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters=64, kernel_size=(5,5),padding='Same',activation=leakyrelu))
model16.add(MaxPooling2D(pool_size=(2,2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters=128, kernel_size=(5,5),padding='Same',activation=leakyrelu))
model16.add(MaxPooling2D(pool_size=(2,2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters=256, kernel_size=(5,5),padding='Same',activation=leakyrelu))
model16.add(MaxPooling2D(pool_size=(2,2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters=512, kernel_size=(5,5),padding='Same',activation=leakyrelu))
model16.add(MaxPooling2D(pool_size=(2,2)))
model16.add(Dropout(0.1))

model16.add(Flatten())
model16.add(Dense(128, activation=leakyrelu))
model16.add(Dense(256, activation=leakyrelu))
model16.add(Dense(5, activation='softmax'))

model16.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model16.summary()
```

Now, we will use this model and run the experiment on color images.

Load the colour dataset

Pre-processing

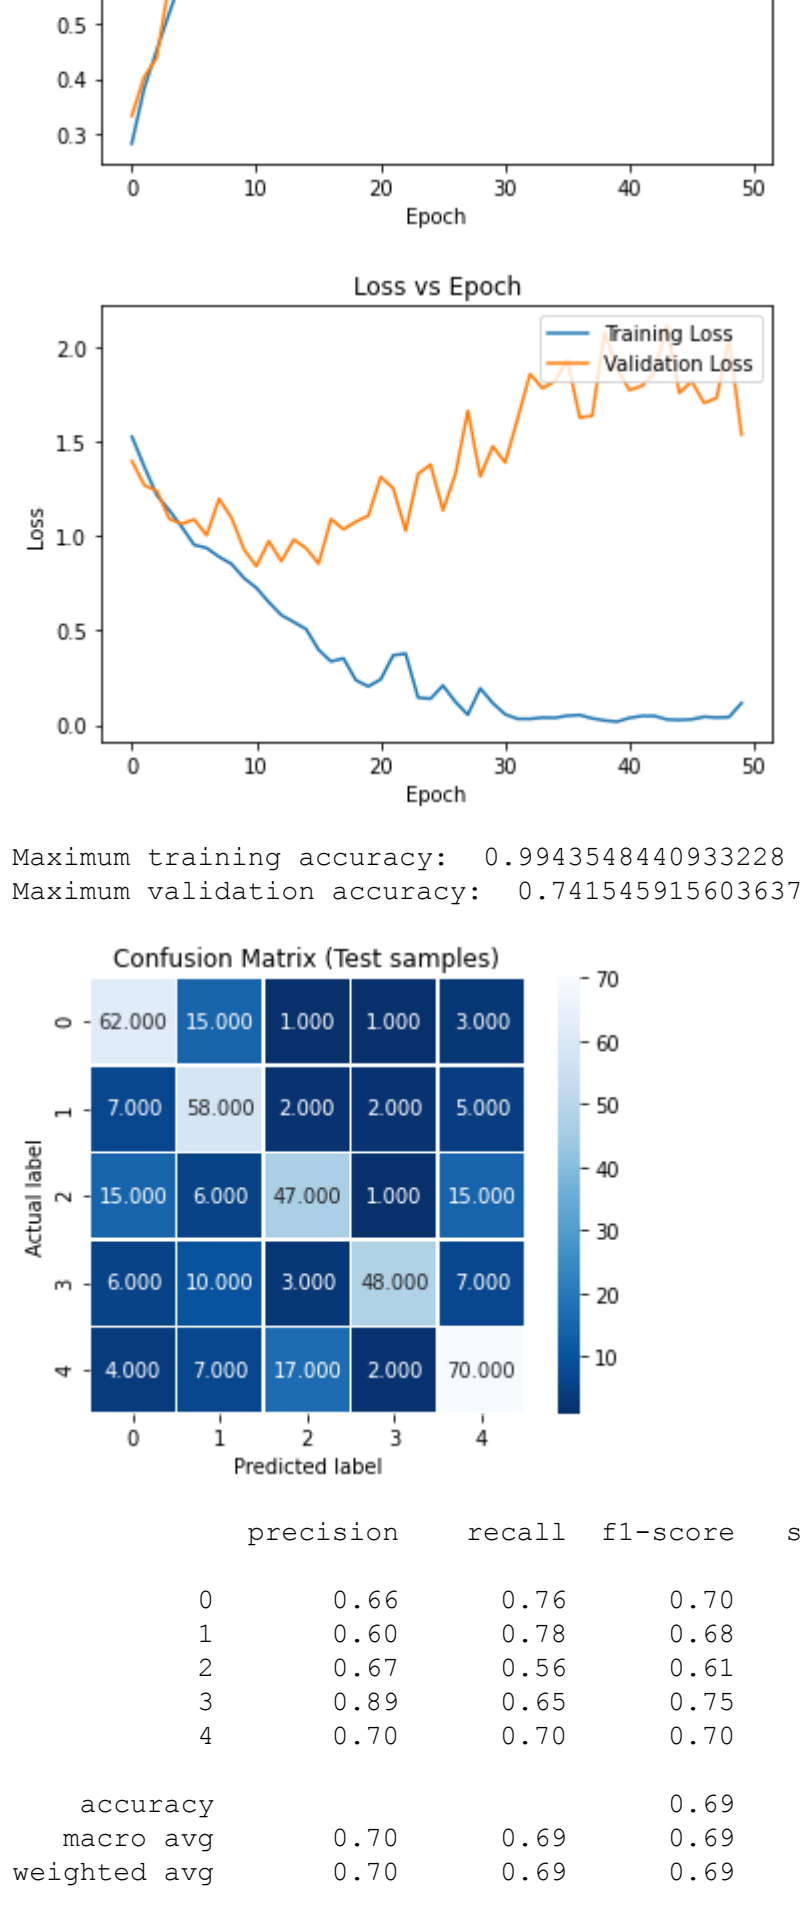
Train : Test split

Training the model

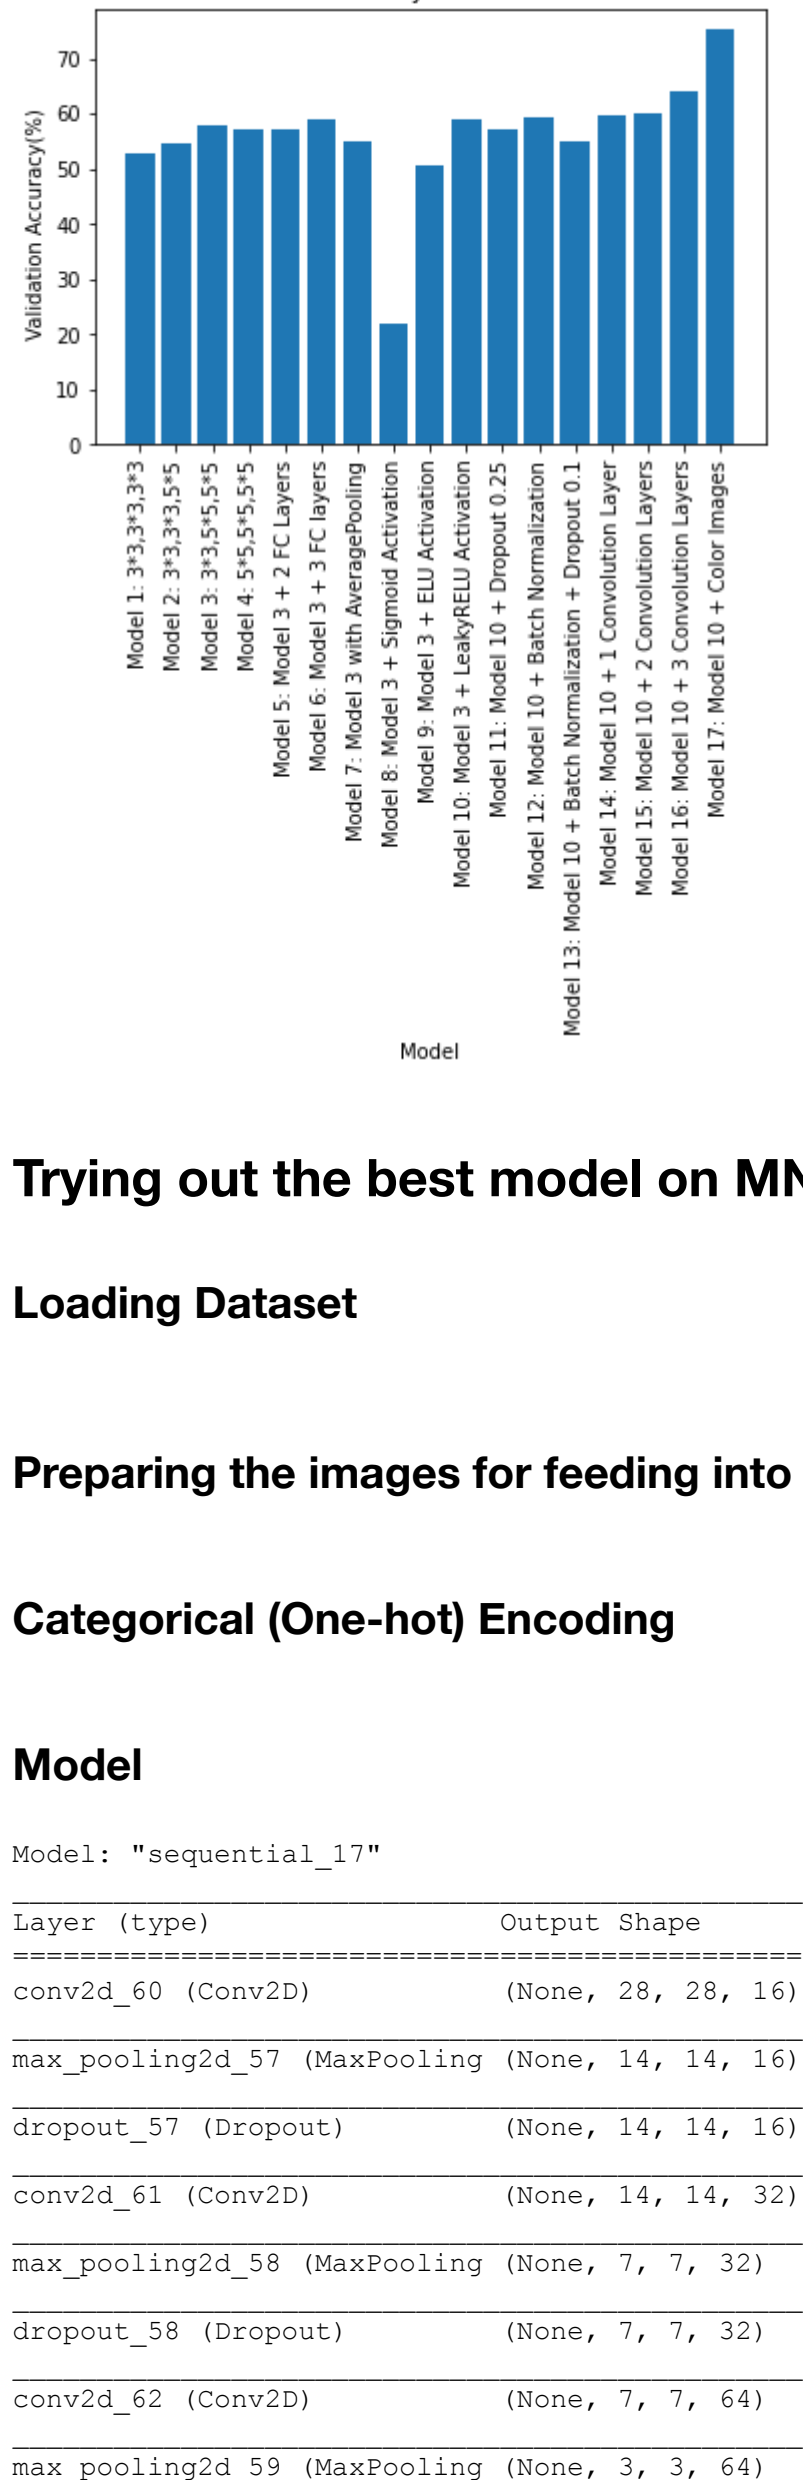
Model: "sequential_16"

Layer (type)	Output Shape	Param #
conv2d_54 (Conv2D)	(None, 80, 80, 16)	448
max_pooling2d_51 (MaxPooling)	(None, 40, 40, 16)	0
dropout_51 (Dropout)	(None, 40, 40, 16)	0
conv2d_55 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_52 (MaxPooling)	(None, 20, 20, 32)	0
dropout_52 (Dropout)	(None, 20, 20, 32)	0
conv2d_56 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_53 (MaxPooling)	(None, 10, 10, 64)	0
dropout_53 (Dropout)	(None, 10, 10, 64)	0
conv2d_57 (Conv2D)	(None, 10, 10, 128)	204928
max_pooling2d_54 (MaxPooling)	(None, 5, 5, 128)	0
dropout_54 (Dropout)	(None, 5, 5, 128)	0
conv2d_58 (Conv2D)	(None, 5, 5, 256)	819456
max_pooling2d_55 (MaxPooling)	(None, 2, 2, 256)	0
dropout_55 (Dropout)	(None, 2, 2, 256)	0
conv2d_59 (Conv2D)	(None, 2, 2, 512)	3277312
max_pooling2d_56 (MaxPooling)	(None, 1, 1, 512)	0
dropout_56 (Dropout)	(None, 1, 1, 512)	0
flatten_16 (Flatten)	(None, 512)	0
dense_39 (Dense)	(None, 128)	65664
dense_40 (Dense)	(None, 256)	33024
dense_41 (Dense)	(None, 5)	1285

=====
Total params: 4,466,213
Trainable params: 4,466,213
Non-trainable params: 0



Plot bar-graph for comparing validation accuracy for all the models



Trying out the best model on MNIST Dataset

Loading Dataset

Preparing the images for feeding into CNN

Categorical (One-hot) Encoding

Model

Model: "sequential_17"


Layer (type)	Output Shape	Param #
conv2d_60 (Conv2D)	(None, 28, 28, 16)	160
max_pooling2d_57 (MaxPooling)	(None, 14, 14, 16)	0
dropout_57 (Dropout)	(None, 14, 14, 16)	0
conv2d_61 (Conv2D)	(None, 14, 14, 32)	12832
max_pooling2d_58 (MaxPooling)	(None, 7, 7, 32)	0
dropout_58 (Dropout)	(None, 7, 7, 32)	0
conv2d_62 (Conv2D)	(None, 7, 7, 64)	51264
max_pooling2d_59 (MaxPooling)	(None, 3, 3, 64)	0
dropout_59 (Dropout)	(None, 3, 3, 64)	0
conv2d_63 (Conv2D)	(None, 3, 3, 128)	204928
max_pooling2d_60 (MaxPooling)	(None, 1, 1, 128)	0
dropout_60 (Dropout)	(None, 1, 1, 128)	0
conv2d_64 (Conv2D)	(None, 1, 1, 256)	819456
dropout_61 (Dropout)	(None, 1, 1, 256)	0
conv2d_65 (Conv2D)	(None, 1, 1, 512)	3277312
dropout_62 (Dropout)	(None, 1, 1, 512)	0
flatten_17 (Flatten)	(None, 512)	0
dense_42 (Dense)	(None, 128)	65664
dense_43 (Dense)	(None, 256)	33024
dense_44 (Dense)	(None, 10)	2570

=====
Total params: 4,467,210
Trainable params: 4,467,210
Non-trainable params: 0


```
Epoch 1/10
422/422 [=====] - 7s 16ms/step - loss: 0.3477 - accuracy: 0.8831 - val_loss: 0.0603 - val_accuracy: 0.9820
Epoch 2/10
422/422 [=====] - 6s 15ms/step - loss: 0.0841 - accuracy: 0.9770 - val_loss: 0.0636 - val_accuracy: 0.9822
Epoch 3/10
422/422 [=====] - 6s 15ms/step - loss: 0.0660 - accuracy: 0.9825 - val_loss: 0.0543 - val_accuracy: 0.9892
Epoch 4/10
422/422 [=====] - 6s 15ms/step - loss: 0.0428 - accuracy: 0.9883 - val_loss: 0.0373 - val_accuracy: 0.9917
Epoch 5/10
422/422 [=====] - 6s 15ms/step - loss: 0.0506 - accuracy: 0.9869 - val_loss: 0.0380 - val_accuracy: 0.9923
Epoch 6/10
422/422 [=====] - 6s 15ms/step - loss: 0.0428 - accuracy: 0.9883 - val_loss: 0.0373 - val_accuracy: 0.9892
Epoch 7/10
422/422 [=====] - 6s 15ms/step - loss: 0.0404 - accuracy: 0.9897 - val_loss: 0.0380 - val_accuracy: 0.9917
Epoch 8/10
422/422 [=====] - 6s 15ms/step - loss: 0.0390 - accuracy: 0.9900 - val_loss: 0.0393 - val_accuracy: 0.9912
Epoch 9/10
422/422 [=====] - 6s 15ms/step - loss: 0.0336 - accuracy: 0.9911 - val_loss: 0.0393 - val_accuracy: 0.9917
Epoch 10/10
422/422 [=====] - 6s 15ms/step - loss: 0.0310 - accuracy: 0.9921 - val_loss: 0.0393 - val_accuracy: 0.9917
```

Making predictions on 5 test images

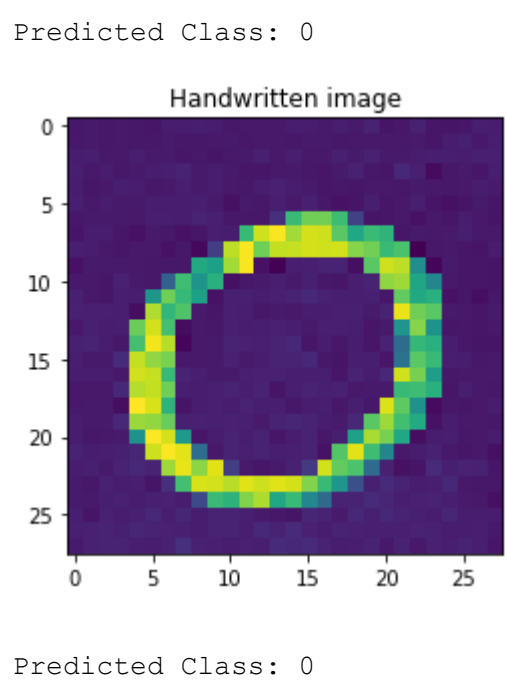
True class: 6
Predicted class: 6




True class: 5
Predicted class: 5




True class: 4
Predicted class: 4



True class: 9
Predicted class: 9




True class: 1
Predicted class: 1



Testing on 5 handwritten images

Predicted Class: 0



Predicted Class: 0

Predicted Class: 3

Predicted Class: 1

Predicted Class: 2

