

Assignment 7

Link to questions - [here](#)

(https://docs.google.com/document/d/1o7tMfmpS5XhLc9YpQHV92ZJs1YgyZ40-G98zOVH_0_8/edit?usp=sharing)

Student Details:

- Name : Anjishnu Mukherjee
- Registration Number : B05-511017020
- Class Roll Number : CS Gy-70
- Exam Roll Number : 510517086
- Email : 511017020.anjishnu@students.iiests.ac.in

Connect to Drive

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Imports ¶

2.3.0

Explore and preprocess

Read CSV

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	H
0	149687	B001EPQRGG	A14OQ67LM79KW3	Candace Mike N Elmo "candacemikenelmo"		0
1	393601	B001ET5XVW	A39CC8SWYBILWO	LHK		1
2	393464	B001SATUF2	A8WIS8WGO3B30	stonrdude "dude"		0
3	479900	B0037QTMD0	A804FIHJ3P9UI	W. J. Costlow		1
4	175089	B002TXT502	A23X9QV9XPU9MG	BekahKnits		1

Assign sentiment and create list for labels & reviews

Convert reviews into sequences of word indices

- A "word index" would be an ID for the word
- Sequences are truncated to a maximum length of 100 words

```
Length of sequences: 393579
Found 133039 unique tokens.
Shape of data tensor: (393579, 100)
Shape of label tensor: (393579, 2)
```

Make separate lists for positive and negative reviews

```
Positive reviews # 86823
Negative reviews # 306756
```

Build train, validation, test data

- 10000 reviews are considered for training
- 2000 reviews are considered for each test and validation sets
- For each of the three segments, equal number of pos. and neg. reviews are taken

```
Training 10000
Validation 2000
Test 2000
```

Storing and retrieving the data variables

Preparing the embedding matrix

- Using the pre-trained GloVe embedding which will contain at index i the embedding vector for the word of index i in our word index.

```
Found 400000 word vectors.
```

Storing and retrieving the data variables

Building Model

Utility function to plot accuracy and loss

Importing keras

GloVe Embedding layer

Model 1 (1 LSTM-64)

WARNING:tensorflow:Layer lstm will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 18s 226ms/step - loss: 0.9867 - accuracy: 0.5639 - val_loss: 0.6866 - val_accuracy: 0.5345

Epoch 2/10

79/79 [=====] - 18s 222ms/step - loss: 0.6711 - accuracy: 0.5991 - val_loss: 0.6743 - val_accuracy: 0.5965

Epoch 3/10

79/79 [=====] - 17s 218ms/step - loss: 0.6672 - accuracy: 0.6130 - val_loss: 0.6726 - val_accuracy: 0.6010

Epoch 4/10

79/79 [=====] - 18s 227ms/step - loss: 0.6650 - accuracy: 0.6139 - val_loss: 0.6710 - val_accuracy: 0.6010

Epoch 5/10

79/79 [=====] - 17s 221ms/step - loss: 0.6627 - accuracy: 0.6191 - val_loss: 0.6697 - val_accuracy: 0.5985

Epoch 6/10

79/79 [=====] - 18s 227ms/step - loss: 0.6603 - accuracy: 0.6202 - val_loss: 0.6672 - val_accuracy: 0.6060

Epoch 7/10

79/79 [=====] - 17s 214ms/step - loss: 0.6232 - accuracy: 0.5924 - val_loss: 0.6931 - val_accuracy: 0.5135

Epoch 8/10

79/79 [=====] - 18s 225ms/step - loss: 0.67365 - accuracy: 0.5675 - val_loss: 0.6709 - val_accuracy: 0.6115

Epoch 9/10

79/79 [=====] - 17s 221ms/step - loss: 0.6673 - accuracy: 0.6115 - val_loss: 0.6694 - val_accuracy: 0.6140

Epoch 10/10



79/79 [=====] - 17s 220ms/step - loss: 0.6662 - accuracy: 0.6120 - val_loss: 0.6683 - val_accuracy: 0.6185

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
lstm (LSTM)	(None, 64)	42240
dense (Dense)	(None, 2)	130
Total params: 13,346,370		
Trainable params: 42,370		
Non-trainable params: 13,304,000		

Evaluate on test data

63/63 [=====] - 1s 19ms/step - loss: 0.6709 - accuracy: 0.5965
test loss, test acc: [0.6708700060844421, 0.5964999794960022]



Model 2 (1 GRU-64)

WARNING:tensorflow:Layer gru will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 21s 267ms/step - loss: 0.6709 - accuracy: 0.5844 - val_loss: 0.6357 - val_accuracy: 0.6450

Epoch 2/10

79/79 [=====] - 23s 290ms/step - loss: 0.5604 - accuracy: 0.7153 - val_loss: 0.4856 - val_accuracy: 0.7705

Epoch 3/10

79/79 [=====] - 22s 284ms/step - loss: 0.4819 - accuracy: 0.7733 - val_loss: 0.4577 - val_accuracy: 0.7820

Epoch 4/10

79/79 [=====] - 23s 285ms/step - loss: 0.4497 - accuracy: 0.7927 - val_loss: 0.4492 - val_accuracy: 0.7970

Epoch 5/10

79/79 [=====] - 22s 280ms/step - loss: 0.4325 - accuracy: 0.8010 - val_loss: 0.4427 - val_accuracy: 0.7975

Epoch 6/10

79/79 [=====] - 21s 266ms/step - loss: 0.4098 - accuracy: 0.8100 - val_loss: 0.4251 - val_accuracy: 0.8085

Epoch 7/10

79/79 [=====] - 22s 278ms/step - loss: 0.3917 - accuracy: 0.8264 - val_loss: 0.4852 - val_accuracy: 0.7735

Epoch 8/10

79/79 [=====] - 22s 274ms/step - loss: 0.3773 - accuracy: 0.8327 - val_loss: 0.4340 - val_accuracy: 0.8065

Epoch 9/10

79/79 [=====] - 21s 268ms/step - loss: 0.3609 - accuracy: 0.8434 - val_loss: 0.4185 - val_accuracy: 0.8110

Epoch 10/10



79/79 [=====] - 21s 264ms/step - loss: 0.3347 - accuracy: 0.8560 - val_loss: 0.4272 - val_accuracy: 0.8115

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru (GRU)	(None, 64)	31872
dense_1 (Dense)	(None, 2)	130
Total params: 13,336,002		
Trainable params: 32,002		
Non-trainable params: 13,304,000		

Evaluate on test data

63/63 [=====] - 1s 21ms/step - loss: 0.39
24 - accuracy: 0.8235
test loss, test acc: [0.39243075251579285, 0.8234999775886536]



Best model so far is model 2 (using GRU)

Model 3 (1 GRU-32)

WARNING:tensorflow:Layer gru_1 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 23s 288ms/step - loss: 0.6789 - accuracy: 0.5679 - val_loss: 0.6570 - val_accuracy: 0.6085

Epoch 2/10

79/79 [=====] - 21s 264ms/step - loss: 0.6003 - accuracy: 0.6712 - val_loss: 0.5199 - val_accuracy: 0.7525

Epoch 3/10

79/79 [=====] - 22s 278ms/step - loss: 0.4948 - accuracy: 0.7585 - val_loss: 0.5435 - val_accuracy: 0.7280

Epoch 4/10

79/79 [=====] - 24s 309ms/step - loss: 0.4648 - accuracy: 0.7767 - val_loss: 0.4513 - val_accuracy: 0.7925

Epoch 5/10

79/79 [=====] - 22s 284ms/step - loss: 0.4419 - accuracy: 0.7918 - val_loss: 0.4550 - val_accuracy: 0.7855

Epoch 6/10

79/79 [=====] - 23s 288ms/step - loss: 0.4283 - accuracy: 0.8016 - val_loss: 0.4308 - val_accuracy: 0.8000

Epoch 7/10

79/79 [=====] - 23s 292ms/step - loss: 0.4088 - accuracy: 0.8109 - val_loss: 0.4280 - val_accuracy: 0.8050

Epoch 8/10

79/79 [=====] - 23s 285ms/step - loss: 0.3985 - accuracy: 0.8190 - val_loss: 0.4185 - val_accuracy: 0.8135

Epoch 9/10

79/79 [=====] - 22s 281ms/step - loss: 0.3829 - accuracy: 0.8237 - val_loss: 0.4172 - val_accuracy: 0.8190

Epoch 10/10



79/79 [=====] - 22s 280ms/step - loss: 0.3769 - accuracy: 0.8314 - val_loss: 0.4167 - val_accuracy: 0.8230

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru_1 (GRU)	(None, 32)	12864
dense_2 (Dense)	(None, 2)	66
Total params: 13,316,930		
Trainable params: 12,930		
Non-trainable params: 13,304,000		

Evaluate on test data

63/63 [=====] - 2s 24ms/step - loss: 0.40
07 - accuracy: 0.8210
test loss, test acc: [0.40069156885147095, 0.8209999799728394]



Model 4 (1 GRU-128)

WARNING:tensorflow:Layer gru_2 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 22s 280ms/step - loss: 0.6614 - accuracy: 0.5955 - val_loss: 0.5921 - val_accuracy: 0.6855

Epoch 2/10

79/79 [=====] - 22s 279ms/step - loss: 0.5088 - accuracy: 0.7517 - val_loss: 0.4978 - val_accuracy: 0.7630

Epoch 3/10

79/79 [=====] - 23s 287ms/step - loss: 0.4607 - accuracy: 0.7873 - val_loss: 0.4864 - val_accuracy: 0.7635

Epoch 4/10

79/79 [=====] - 23s 288ms/step - loss: 0.4268 - accuracy: 0.8046 - val_loss: 0.4454 - val_accuracy: 0.7925

Epoch 5/10

79/79 [=====] - 22s 272ms/step - loss: 0.3984 - accuracy: 0.8185 - val_loss: 0.4240 - val_accuracy: 0.8090

Epoch 6/10

79/79 [=====] - 22s 284ms/step - loss: 0.3747 - accuracy: 0.8378 - val_loss: 0.4358 - val_accuracy: 0.7995

Epoch 7/10

79/79 [=====] - 22s 282ms/step - loss: 0.3479 - accuracy: 0.8439 - val_loss: 0.4010 - val_accuracy: 0.8200

Epoch 8/10

79/79 [=====] - 22s 273ms/step - loss: 0.3173 - accuracy: 0.8635 - val_loss: 0.5076 - val_accuracy: 0.7775

Epoch 9/10

79/79 [=====] - 23s 288ms/step - loss: 0.3111 - accuracy: 0.8695 - val_loss: 0.4352 - val_accuracy: 0.8055

Epoch 10/10

79/79 [=====] - 22s 284ms/step - loss: 0.2866 - accuracy: 0.8809 - val_loss: 0.4263 - val_accuracy: 0.8150

Model: "sequential_3"



Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru_2 (GRU)	(None, 128)	88320
dense_3 (Dense)	(None, 2)	258
Total params: 13,392,578		
Trainable params: 88,578		
Non-trainable params: 13,304,000		

Evaluate on test data

63/63 [=====] - 2s 25ms/step - loss: 0.38

92 - accuracy: 0.8385

test loss, test acc: [0.389176607131958, 0.8385000228881836]



Now, Model 4 is the best performing

For stacked LSTMs/GRUs, add `return_sequences=True` for all but the last of those layers

Model 5 (2 GRU-128)

```
WARNING:tensorflow:Layer gru_3 will not use cuDNN kernel since it
doesn't meet the cuDNN kernel criteria. It will use generic GPU ke
rnel as fallback when running on GPU
```

```
WARNING:tensorflow:Layer gru_4 will not use cuDNN kernel since it
doesn't meet the cuDNN kernel criteria. It will use generic GPU ke
rnel as fallback when running on GPU
```

```
Epoch 1/10
```

```
79/79 [=====] - 46s 576ms/step - loss: 0.
6601 - accuracy: 0.5907 - val_loss: 0.6114 - val_accuracy: 0.6945
```

```
Epoch 2/10
```

```
79/79 [=====] - 45s 565ms/step - loss: 0.
5187 - accuracy: 0.7409 - val_loss: 0.4645 - val_accuracy: 0.7810
```

```
Epoch 3/10
```

```
79/79 [=====] - 43s 545ms/step - loss: 0.
4365 - accuracy: 0.8002 - val_loss: 0.4318 - val_accuracy: 0.8020
```

```
Epoch 4/10
```

```
79/79 [=====] - 44s 557ms/step - loss: 0.
4137 - accuracy: 0.8116 - val_loss: 0.4131 - val_accuracy: 0.8155
```

```
Epoch 5/10
```

```
79/79 [=====] - 45s 567ms/step - loss: 0.
3721 - accuracy: 0.8365 - val_loss: 0.4069 - val_accuracy: 0.8180
```

```
Epoch 6/10
```

```
79/79 [=====] - 44s 556ms/step - loss: 0.
3492 - accuracy: 0.8487 - val_loss: 0.3951 - val_accuracy: 0.8275
```

```
Epoch 7/10
```

```
79/79 [=====] - 45s 571ms/step - loss: 0.
3226 - accuracy: 0.8611 - val_loss: 0.4206 - val_accuracy: 0.8185
```

```
Epoch 8/10
```

```
79/79 [=====] - 45s 564ms/step - loss: 0.
2951 - accuracy: 0.8772 - val_loss: 0.4137 - val_accuracy: 0.8240
```

```
Epoch 9/10
```

```
79/79 [=====] - 45s 573ms/step - loss: 0.
2571 - accuracy: 0.8968 - val_loss: 0.4514 - val_accuracy: 0.8135
```

```
Epoch 10/10
```

```
79/79 [=====] - 44s 562ms/step - loss: 0.
2481 - accuracy: 0.8986 - val_loss: 0.4286 - val_accuracy: 0.8200
```

Model: "sequential_4"


Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru_3 (GRU)	(None, 100, 128)	88320
gru_4 (GRU)	(None, 128)	99072
dense_4 (Dense)	(None, 2)	258
Total params: 13,491,650		
Trainable params: 187,650		
Non-trainable params: 13,304,000		


Evaluate on test data

63/63 [=====] - 2s 40ms/step - loss: 0.40

28 - accuracy: 0.8375

test loss, test acc: [0.4027763307094574, 0.8374999761581421]





WARNING:tensorflow:Layer gru_3 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_4 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

63/63 [=====] - 3s 47ms/step - loss: 0.4028 - accuracy: 0.8375

Model 6 (3 GRU-128)

WARNING:tensorflow:Layer gru_5 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_6 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_7 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 68s 856ms/step - loss: 0.6592 - accuracy: 0.5894 - val_loss: 0.5878 - val_accuracy: 0.6900

Epoch 2/10

79/79 [=====] - 66s 830ms/step - loss: 0.5170 - accuracy: 0.7471 - val_loss: 0.5811 - val_accuracy: 0.7065

Epoch 3/10

79/79 [=====] - 65s 818ms/step - loss: 0.4507 - accuracy: 0.7899 - val_loss: 0.4405 - val_accuracy: 0.7935

Epoch 4/10

79/79 [=====] - 65s 822ms/step - loss: 0.4192 - accuracy: 0.8047 - val_loss: 0.4772 - val_accuracy: 0.7740

Epoch 5/10

79/79 [=====] - 65s 829ms/step - loss: 0.3786 - accuracy: 0.8284 - val_loss: 0.4281 - val_accuracy: 0.8015

Epoch 6/10

79/79 [=====] - 65s 826ms/step - loss: 0.3764 - accuracy: 0.8293 - val_loss: 0.3978 - val_accuracy: 0.8240

Epoch 7/10

79/79 [=====] - 65s 825ms/step - loss: 0.3368 - accuracy: 0.8531 - val_loss: 0.3952 - val_accuracy: 0.8315

Epoch 8/10

79/79 [=====] - 64s 813ms/step - loss: 0.3011 - accuracy: 0.8740 - val_loss: 0.5142 - val_accuracy: 0.7865

Epoch 9/10

79/79 [=====] - 66s 835ms/step - loss: 0.2738 - accuracy: 0.8855 - val_loss: 0.4639 - val_accuracy: 0.8055

Epoch 10/10

79/79 [=====] - 65s 819ms/step - loss: 0.2495 - accuracy: 0.8970 - val_loss: 0.4446 - val_accuracy: 0.8215

Model: "sequential_5"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru_5 (GRU)	(None, 100, 128)	88320
gru_6 (GRU)	(None, 100, 128)	99072
gru_7 (GRU)	(None, 128)	99072
dense_5 (Dense)	(None, 2)	258

Total params: 13,590,722

Trainable params: 286,722


Non-trainable params: 13,304,000


Evaluate on test data

63/63 [=====] - 4s 60ms/step - loss: 0.40

78 - accuracy: 0.8320

test loss, test acc: [0.40783989429473877, 0.8320000171661377]





Model 7 (4 GRU-128)

WARNING:tensorflow:Layer gru_8 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_9 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_10 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_11 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 85s 1s/step - loss: 0.6706 - accuracy: 0.5734 - val_loss: 0.6438 - val_accuracy: 0.6495

Epoch 2/10

79/79 [=====] - 84s 1s/step - loss: 0.5240 - accuracy: 0.7448 - val_loss: 0.4785 - val_accuracy: 0.7750

Epoch 3/10

79/79 [=====] - 84s 1s/step - loss: 0.4418 - accuracy: 0.7947 - val_loss: 0.4533 - val_accuracy: 0.7855

Epoch 4/10

79/79 [=====] - 85s 1s/step - loss: 0.4117 - accuracy: 0.8136 - val_loss: 0.4018 - val_accuracy: 0.8150

Epoch 5/10

79/79 [=====] - 85s 1s/step - loss: 0.3696 - accuracy: 0.8361 - val_loss: 0.4564 - val_accuracy: 0.7865

Epoch 6/10

79/79 [=====] - 84s 1s/step - loss: 0.3660 - accuracy: 0.8414 - val_loss: 0.4083 - val_accuracy: 0.8210

Epoch 7/10

79/79 [=====] - 84s 1s/step - loss: 0.3218 - accuracy: 0.8664 - val_loss: 0.4090 - val_accuracy: 0.8165

Epoch 8/10

79/79 [=====] - 86s 1s/step - loss: 0.2978 - accuracy: 0.8747 - val_loss: 0.3936 - val_accuracy: 0.8235

Epoch 9/10

79/79 [=====] - 83s 1s/step - loss: 0.2618 - accuracy: 0.8968 - val_loss: 0.4632 - val_accuracy: 0.7965

Epoch 10/10

79/79 [=====] - 84s 1s/step - loss: 0.2599 - accuracy: 0.8940 - val_loss: 0.4468 - val_accuracy: 0.8060

Model: "sequential_6"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru_8 (GRU)	(None, 100, 128)	88320
gru_9 (GRU)	(None, 100, 128)	99072
gru_10 (GRU)	(None, 100, 128)	99072
gru_11 (GRU)	(None, 128)	99072
dense_6 (Dense)	(None, 2)	258

Total params: 13,689,794

Trainable params: 385,794


Non-trainable params: 13,304,000


Evaluate on test data

63/63 [=====] - 4s 70ms/step - loss: 0.40

19 - accuracy: 0.8260

test loss, test acc: [0.4018857479095459, 0.8259999752044678]





Model 5 is the best performing so far

Model 8 (2 Bidirectional GRU-128)

```
WARNING:tensorflow:Layer gru_12 will not use cuDNN kernel since it
doesn't meet the cuDNN kernel criteria. It will use generic GPU ke
rnel as fallback when running on GPU
WARNING:tensorflow:Layer gru_12 will not use cuDNN kernel since it
doesn't meet the cuDNN kernel criteria. It will use generic GPU ke
rnel as fallback when running on GPU
WARNING:tensorflow:Layer gru_12 will not use cuDNN kernel since it
doesn't meet the cuDNN kernel criteria. It will use generic GPU ke
rnel as fallback when running on GPU
WARNING:tensorflow:Layer gru_13 will not use cuDNN kernel since it
doesn't meet the cuDNN kernel criteria. It will use generic GPU ke
rnel as fallback when running on GPU
WARNING:tensorflow:Layer gru_13 will not use cuDNN kernel since it
doesn't meet the cuDNN kernel criteria. It will use generic GPU ke
rnel as fallback when running on GPU
WARNING:tensorflow:Layer gru_13 will not use cuDNN kernel since it
doesn't meet the cuDNN kernel criteria. It will use generic GPU ke
rnel as fallback when running on GPU
Epoch 1/10
79/79 [=====] - 83s 1s/step - loss: 0.640
7 - accuracy: 0.6191 - val_loss: 0.6861 - val_accuracy: 0.5845
Epoch 2/10
79/79 [=====] - 85s 1s/step - loss: 0.499
8 - accuracy: 0.7569 - val_loss: 0.6557 - val_accuracy: 0.6735
Epoch 3/10
79/79 [=====] - 83s 1s/step - loss: 0.466
0 - accuracy: 0.7813 - val_loss: 0.4682 - val_accuracy: 0.7725
Epoch 4/10
79/79 [=====] - 83s 1s/step - loss: 0.394
9 - accuracy: 0.8208 - val_loss: 0.4095 - val_accuracy: 0.8170
Epoch 5/10
79/79 [=====] - 82s 1s/step - loss: 0.364
8 - accuracy: 0.8387 - val_loss: 0.4522 - val_accuracy: 0.7875
Epoch 6/10
79/79 [=====] - 80s 1s/step - loss: 0.353
4 - accuracy: 0.8446 - val_loss: 0.4024 - val_accuracy: 0.8295
Epoch 7/10
79/79 [=====] - 82s 1s/step - loss: 0.300
1 - accuracy: 0.8729 - val_loss: 0.4158 - val_accuracy: 0.8195
Epoch 8/10
79/79 [=====] - 81s 1s/step - loss: 0.269
9 - accuracy: 0.8889 - val_loss: 0.4029 - val_accuracy: 0.8365
Epoch 9/10
79/79 [=====] - 82s 1s/step - loss: 0.228
8 - accuracy: 0.9097 - val_loss: 0.4358 - val_accuracy: 0.8310
Epoch 10/10
79/79 [=====] - 82s 1s/step - loss: 0.196
3 - accuracy: 0.9238 - val_loss: 0.5506 - val_accuracy: 0.7740
```

Model: "sequential_7"


Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
bidirectional (Bidirectional)	(None, 100, 256)	176640
bidirectional_1 (Bidirectional)	(None, 256)	296448
dense_7 (Dense)	(None, 2)	514
Total params: 13,777,602		
Trainable params: 473,602		
Non-trainable params: 13,304,000		


Evaluate on test data

63/63 [=====] - 5s 72ms/step - loss: 0.54

24 - accuracy: 0.7910

test loss, test acc: [0.5423805117607117, 0.7910000085830688]





Still, Model 5 is the best performing

Model 9 (2 GRU-128 dropout=0.1)

WARNING:tensorflow:Layer gru_14 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_15 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 48s 604ms/step - loss: 0.6687 - accuracy: 0.5840 - val_loss: 0.6305 - val_accuracy: 0.6575

Epoch 2/10

79/79 [=====] - 47s 599ms/step - loss: 0.5712 - accuracy: 0.7052 - val_loss: 0.5294 - val_accuracy: 0.7365

Epoch 3/10

79/79 [=====] - 49s 622ms/step - loss: 0.4826 - accuracy: 0.7644 - val_loss: 0.4429 - val_accuracy: 0.7920

Epoch 4/10

79/79 [=====] - 48s 609ms/step - loss: 0.4367 - accuracy: 0.7960 - val_loss: 0.4267 - val_accuracy: 0.8070

Epoch 5/10

79/79 [=====] - 49s 619ms/step - loss: 0.4190 - accuracy: 0.8072 - val_loss: 0.4235 - val_accuracy: 0.8085

Epoch 6/10

79/79 [=====] - 49s 614ms/step - loss: 0.3955 - accuracy: 0.8219 - val_loss: 0.4039 - val_accuracy: 0.8175

Epoch 7/10

79/79 [=====] - 45s 574ms/step - loss: 0.3775 - accuracy: 0.8313 - val_loss: 0.4066 - val_accuracy: 0.8210

Epoch 8/10

79/79 [=====] - 47s 595ms/step - loss: 0.3568 - accuracy: 0.8470 - val_loss: 0.3973 - val_accuracy: 0.8285

Epoch 9/10

79/79 [=====] - 47s 595ms/step - loss: 0.3279 - accuracy: 0.8607 - val_loss: 0.4001 - val_accuracy: 0.8260

Epoch 10/10

79/79 [=====] - 48s 604ms/step - loss: 0.3099 - accuracy: 0.8686 - val_loss: 0.4594 - val_accuracy: 0.8110

Model: "sequential_8"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru_14 (GRU)	(None, 100, 128)	88320
gru_15 (GRU)	(None, 128)	99072
dense_8 (Dense)	(None, 2)	258

Total params: 13,491,650

Trainable params: 187,650


Non-trainable params: 13,304,000


Evaluate on test data

63/63 [=====] - 2s 37ms/step - loss: 0.42

05 - accuracy: 0.8200

test loss, test acc: [0.4205341935157776, 0.8199999928474426]





Model 10 (2 GRU-128 recurrent dropout=0.2)

WARNING:tensorflow:Layer gru will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_1 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 43s 540ms/step - loss: 0.6658 - accuracy: 0.5871 - val_loss: 0.6432 - val_accuracy: 0.6335

Epoch 2/10

79/79 [=====] - 42s 531ms/step - loss: 0.5327 - accuracy: 0.7316 - val_loss: 0.4648 - val_accuracy: 0.7860

Epoch 3/10

79/79 [=====] - 42s 532ms/step - loss: 0.4579 - accuracy: 0.7850 - val_loss: 0.5071 - val_accuracy: 0.7505

Epoch 4/10

79/79 [=====] - 42s 533ms/step - loss: 0.4273 - accuracy: 0.8065 - val_loss: 0.4267 - val_accuracy: 0.8005

Epoch 5/10

79/79 [=====] - 42s 532ms/step - loss: 0.3967 - accuracy: 0.8218 - val_loss: 0.4169 - val_accuracy: 0.8120

Epoch 6/10

79/79 [=====] - 42s 527ms/step - loss: 0.3631 - accuracy: 0.8436 - val_loss: 0.5645 - val_accuracy: 0.7385

Epoch 7/10

79/79 [=====] - 42s 527ms/step - loss: 0.3517 - accuracy: 0.8506 - val_loss: 0.3982 - val_accuracy: 0.8220

Epoch 8/10

79/79 [=====] - 42s 529ms/step - loss: 0.3270 - accuracy: 0.8579 - val_loss: 0.3995 - val_accuracy: 0.8250

Epoch 9/10

79/79 [=====] - 42s 530ms/step - loss: 0.2981 - accuracy: 0.8744 - val_loss: 0.4113 - val_accuracy: 0.8285

Epoch 10/10

79/79 [=====] - 42s 526ms/step - loss: 0.2735 - accuracy: 0.8862 - val_loss: 0.4144 - val_accuracy: 0.8325

Model: "sequential"


Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru (GRU)	(None, 100, 128)	88320
gru_1 (GRU)	(None, 128)	99072
dense (Dense)	(None, 2)	258
Total params: 13,491,650		
Trainable params: 187,650		
Non-trainable params: 13,304,000		


Evaluate on test data

63/63 [=====] - 3s 48ms/step - loss: 0.38

61 - accuracy: 0.8365

test loss, test acc: [0.38611656427383423, 0.8364999890327454]





Model 11 (2 GRU-128 dropout=0.1 + recurrent dropout=0.2)

WARNING:tensorflow:Layer gru_2 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_3 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 49s 617ms/step - loss: 0.6742 - accuracy: 0.5827 - val_loss: 0.6221 - val_accuracy: 0.6575

Epoch 2/10

79/79 [=====] - 48s 605ms/step - loss: 0.5510 - accuracy: 0.7198 - val_loss: 0.4776 - val_accuracy: 0.7770

Epoch 3/10

79/79 [=====] - 48s 612ms/step - loss: 0.4656 - accuracy: 0.7804 - val_loss: 0.4649 - val_accuracy: 0.7755

Epoch 4/10

79/79 [=====] - 47s 599ms/step - loss: 0.4531 - accuracy: 0.7871 - val_loss: 0.4272 - val_accuracy: 0.8040

Epoch 5/10

79/79 [=====] - 47s 595ms/step - loss: 0.4138 - accuracy: 0.8073 - val_loss: 0.4218 - val_accuracy: 0.8080

Epoch 6/10

79/79 [=====] - 47s 596ms/step - loss: 0.3910 - accuracy: 0.8233 - val_loss: 0.3976 - val_accuracy: 0.8165

Epoch 7/10

79/79 [=====] - 47s 592ms/step - loss: 0.3662 - accuracy: 0.8372 - val_loss: 0.5794 - val_accuracy: 0.7475

Epoch 8/10

79/79 [=====] - 47s 593ms/step - loss: 0.3729 - accuracy: 0.8344 - val_loss: 0.4045 - val_accuracy: 0.8185

Epoch 9/10

79/79 [=====] - 47s 594ms/step - loss: 0.3385 - accuracy: 0.8535 - val_loss: 0.3932 - val_accuracy: 0.8285

Epoch 10/10


79/79 [=====] - 47s 599ms/step - loss: 0.3212 - accuracy: 0.8627 - val_loss: 0.3951 - val_accuracy: 0.8335


Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru_2 (GRU)	(None, 100, 128)	88320
gru_3 (GRU)	(None, 128)	99072
dense_1 (Dense)	(None, 2)	258
Total params: 13,491,650		
Trainable params: 187,650		
Non-trainable params: 13,304,000		

Evaluate on test data

63/63 [=====] - 3s 46ms/step - loss: 0.3601 - accuracy: 0.8440
test loss, test acc: [0.3601447641849518, 0.843999981880188]





Model 12 (2 GRU-128 bias_regularizer=0.05)

WARNING:tensorflow:Layer gru_4 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_5 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 28s 353ms/step - loss: 0.6606 - accuracy: 0.6134 - val_loss: 0.5524 - val_accuracy: 0.7325

Epoch 2/10

79/79 [=====] - 27s 341ms/step - loss: 0.5203 - accuracy: 0.7509 - val_loss: 0.5140 - val_accuracy: 0.7485

Epoch 3/10

79/79 [=====] - 27s 343ms/step - loss: 0.4593 - accuracy: 0.7901 - val_loss: 0.4940 - val_accuracy: 0.7730

Epoch 4/10

79/79 [=====] - 27s 347ms/step - loss: 0.4259 - accuracy: 0.8073 - val_loss: 0.4768 - val_accuracy: 0.7810

Epoch 5/10

79/79 [=====] - 27s 343ms/step - loss: 0.3972 - accuracy: 0.8252 - val_loss: 0.4185 - val_accuracy: 0.8085

Epoch 6/10

79/79 [=====] - 27s 339ms/step - loss: 0.3745 - accuracy: 0.8391 - val_loss: 0.4047 - val_accuracy: 0.8220

Epoch 7/10

79/79 [=====] - 27s 346ms/step - loss: 0.3448 - accuracy: 0.8567 - val_loss: 0.4238 - val_accuracy: 0.8205

Epoch 8/10

79/79 [=====] - 27s 340ms/step - loss: 0.3200 - accuracy: 0.8686 - val_loss: 0.4423 - val_accuracy: 0.8005

Epoch 9/10

79/79 [=====] - 27s 344ms/step - loss: 0.3149 - accuracy: 0.8678 - val_loss: 0.4097 - val_accuracy: 0.8275

Epoch 10/10


79/79 [=====] - 28s 349ms/step - loss: 0.2732 - accuracy: 0.8908 - val_loss: 0.4286 - val_accuracy: 0.8225


Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	13304000
gru_4 (GRU)	(None, 100, 128)	88320
gru_5 (GRU)	(None, 128)	99072
dense_2 (Dense)	(None, 2)	258
Total params: 13,491,650		
Trainable params: 187,650		
Non-trainable params: 13,304,000		

Evaluate on test data

63/63 [=====] - 2s 32ms/step - loss: 0.4138 - accuracy: 0.8295
test loss, test acc: [0.4138040244579315, 0.8295000195503235]





WARNING:tensorflow:Layer gru_4 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_5 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

63/63 [=====] - 2s 32ms/step - loss: 0.4138 - accuracy: 0.8295

Now, Model 12 is the best performing

Model 13 (2 GRU-128 bias_regularizer=0.05 self-trainable embedding)

WARNING:tensorflow:Layer gru_6 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_7 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/10

79/79 [=====] - 39s 488ms/step - loss: 0.6252 - accuracy: 0.6569 - val_loss: 0.5707 - val_accuracy: 0.7460

Epoch 2/10

79/79 [=====] - 38s 477ms/step - loss: 0.3760 - accuracy: 0.8526 - val_loss: 0.4734 - val_accuracy: 0.7925

Epoch 3/10

79/79 [=====] - 37s 472ms/step - loss: 0.2325 - accuracy: 0.9137 - val_loss: 0.5393 - val_accuracy: 0.7915

Epoch 4/10

79/79 [=====] - 38s 475ms/step - loss: 0.1441 - accuracy: 0.9496 - val_loss: 0.6843 - val_accuracy: 0.7805

Epoch 5/10

79/79 [=====] - 37s 465ms/step - loss: 0.1200 - accuracy: 0.9617 - val_loss: 0.7644 - val_accuracy: 0.7735

Epoch 6/10

79/79 [=====] - 37s 465ms/step - loss: 0.0667 - accuracy: 0.9807 - val_loss: 0.9525 - val_accuracy: 0.7650

Epoch 7/10

79/79 [=====] - 37s 463ms/step - loss: 0.0698 - accuracy: 0.9778 - val_loss: 1.0368 - val_accuracy: 0.7765

Epoch 8/10

79/79 [=====] - 36s 462ms/step - loss: 0.0384 - accuracy: 0.9897 - val_loss: 1.2644 - val_accuracy: 0.7750

Epoch 9/10

79/79 [=====] - 36s 458ms/step - loss: 0.0277 - accuracy: 0.9946 - val_loss: 1.6846 - val_accuracy: 0.7730

Epoch 10/10


79/79 [=====] - 36s 459ms/step - loss: 0.1651 - accuracy: 0.9839 - val_loss: 1.0419 - val_accuracy: 0.7410


Model: "sequential_3"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 100, 100)	13304000
gru_6 (GRU)	(None, 100, 128)	88320
gru_7 (GRU)	(None, 128)	99072
dense_3 (Dense)	(None, 2)	258
Total params: 13,491,650		
Trainable params: 13,491,650		
Non-trainable params: 0		

Evaluate on test data

63/63 [=====] - 2s 31ms/step - loss: 0.9908 - accuracy: 0.7665
test loss, test acc: [0.9908018708229065, 0.7664999961853027]





Model 14 (2 GRU-128 bias_regularizer=0.05 one hot encoded embedding)

- One hot encoding the labels crashes RAM. This is a problem I dont know how to fix. Hence currently not training this model.

Checking on my own review

Length of sequences: 2
Found 23 unique tokens.

```
array([[0.08688124, 0.9131187 ],  
       [0.71508634, 0.2849137 ]], dtype=float32)
```

Hindi movie reviews dataset

Read CSV

	text	experience
0	\n\nबैनर : \nफॉक्स स्टार स्टुडियो, वाइड फ्रेम प...	2
1	Chandermohan.sharma@timesgroup.com बॉलिवुड मे...	0
2	\nडर @ द माल का सबसे बड़ा आकर्षण निर्देशक पवन ...	0
3	सुभाष नागरे या सरकार (अमिताभ बच्चन) महाराष्ट्र...	2
4	बैनर : \nपीवीआर पिक्चर्स\n\nनिर्माता : \nअजय बिज...	2

	text	experience
0	\n\nफिल्म के हीरो संजू (इमरान) और उसके दो साथि...	1
1	\nविक्रम सेठी (सिद्धार्थ मल्होत्रा) और माया (...)	1
2	बैनर : \nयूटीवी मोशन पिक्चर्स, आमिर खान प्रोडक्...	2
3	बैनर : \nयूटीवी मोशन पिक्चर्स, भंडारकर एंटरटेनम...	1
4	\nनिर्माता-निर्देशक विशाल भारद्वाज को साहित्य ...	2

Assign sentiment and create list for labels & reviews

Convert reviews into sequences of word indices

- A "word index" would be an ID for the word
- Sequences are truncated to a maximum length of 100 words

```
Length of sequences: 898
Found 26030 unique tokens.
Shape of data tensor: (898, 100)
Shape of label tensor: (898, 3)
```

Build train, test data

- 80% for training, 10% for validation and 10% for test

```
Training 710
Validation 94
Test 94
```

Model X (2 GRU-128 bias_regularizer=0.05 self-trainable embedding)

WARNING:tensorflow:Layer gru_8 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer gru_9 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Epoch 1/20

23/23 [=====] - 9s 386ms/step - loss: 1.1199 - accuracy: 0.3634 - val_loss: 1.1104 - val_accuracy: 0.3936

Epoch 2/20

23/23 [=====] - 8s 364ms/step - loss: 1.0910 - accuracy: 0.3859 - val_loss: 1.1136 - val_accuracy: 0.3936

Epoch 3/20

23/23 [=====] - 8s 369ms/step - loss: 22.1958 - accuracy: 0.5141 - val_loss: 1.2796 - val_accuracy: 0.3617

Epoch 4/20

23/23 [=====] - 8s 360ms/step - loss: 1.2270 - accuracy: 0.7169 - val_loss: 1.3538 - val_accuracy: 0.3404

Epoch 5/20

23/23 [=====] - 8s 368ms/step - loss: 1.2641 - accuracy: 0.7394 - val_loss: 1.4439 - val_accuracy: 0.4787

Epoch 6/20

23/23 [=====] - 8s 361ms/step - loss: 1.3546 - accuracy: 0.5831 - val_loss: 1.4730 - val_accuracy: 0.4574

Epoch 7/20

23/23 [=====] - 8s 369ms/step - loss: 1.2465 - accuracy: 0.5803 - val_loss: 4.4566 - val_accuracy: 0.4574

Epoch 8/20

23/23 [=====] - 8s 368ms/step - loss: 1.1914 - accuracy: 0.8028 - val_loss: 1.4209 - val_accuracy: 0.3936

Epoch 9/20

23/23 [=====] - 9s 378ms/step - loss: 0.9720 - accuracy: 0.8394 - val_loss: 14612844.0000 - val_accuracy: 0.4043

Epoch 10/20

23/23 [=====] - 9s 374ms/step - loss: 6.9395 - accuracy: 0.5592 - val_loss: 1.8813 - val_accuracy: 0.3511

Epoch 11/20

23/23 [=====] - 9s 374ms/step - loss: 1.4071 - accuracy: 0.5873 - val_loss: 1.8195 - val_accuracy: 0.3298

Epoch 12/20

23/23 [=====] - 8s 367ms/step - loss: 1.3365 - accuracy: 0.6324 - val_loss: 1.7965 - val_accuracy: 0.3404

Epoch 13/20

23/23 [=====] - 8s 359ms/step - loss: 1.2526 - accuracy: 0.6634 - val_loss: 1.7717 - val_accuracy: 0.3617

Epoch 14/20

23/23 [=====] - 8s 362ms/step - loss: 1.1332 - accuracy: 0.7465 - val_loss: 1.6641 - val_accuracy: 0.4255

Epoch 15/20

23/23 [=====] - 8s 367ms/step - loss: 7.4908 - accuracy: 0.8338 - val_loss: 1.7626 - val_accuracy: 0.3830


```

Epoch 16/20
23/23 [=====] - 9s 386ms/step - loss: 1.2
299 - accuracy: 0.7423 - val_loss: 1.8922 - val_accuracy: 0.3191
Epoch 17/20
23/23 [=====] - 8s 368ms/step - loss: 1.2
620 - accuracy: 0.7296 - val_loss: 1.8829 - val_accuracy: 0.3298
Epoch 18/20
23/23 [=====] - 9s 371ms/step - loss: 1.4
068 - accuracy: 0.8070 - val_loss: 1.6987 - val_accuracy: 0.4255
Epoch 19/20
23/23 [=====] - 9s 373ms/step - loss: 1.2
571 - accuracy: 0.7690 - val_loss: 1.6538 - val_accuracy: 0.4362
Epoch 20/20
23/23 [=====] - 8s 367ms/step - loss: 1.3
676 - accuracy: 0.8465 - val_loss: 1.6537 - val_accuracy: 0.4255

```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 100, 100)	2603100
gru_8 (GRU)	(None, 100, 128)	88320
gru_9 (GRU)	(None, 128)	99072
dense_4 (Dense)	(None, 3)	387
Total params: 2,790,879		
Trainable params: 2,790,879		
Non-trainable params: 0		

Evaluate on test data

```

3/3 [=====] - 0s 24ms/step - loss: 1.6553
- accuracy: 0.4149
test loss, test acc: [1.6553058624267578, 0.41489362716674805]

```

